

CS5223 Distributed Systems

Lecture 5: Synchronization (Time and Clock)

Instructor: YU Haifeng

Today's Roadmap

- Chapter 5 of textbook
- Physical clock synchronization
- Software clocks and applications
- Mutual exclusion
- Leader election
- Some contents (e.g., logical clocks and happened-before relation) overlap with CS4231
 - These concepts are critical in distributed system
 - So we need to cover these

Motivation for Physical Clocks

- Example: Timestamp files
 - Creation time, modification time, etc.
- Example: cron jobs
 - Virus scan
 - Software update
- Advantage of physical clock reading: **Portable and context-free**
 - Reading in one system will make sense even in other context

Motivation for Physical Clock Synchronization

- Electronic devices (including computers) usually advances clocks based on oscillations from some crystal
- **Clock rate drift**: Caused by slight differences in oscillation frequency
 - E.g., 1 minute my time = 2 minute your time
- **Clock drift**: Caused by clock rate drift and also lack of synchronization
 - E.g., your clock is 10 minute ahead of the correct time
- **Physical clock synchronization usually only deal with clock drift**

Clock Synchronization Can be Confusing!

- Human beings are usually used to talk about a single clock
 - Confusion usually arises when there are multiple clocks each may advance at its own pace
 - Such confusion already arises in some places in the textbook...
- Need a crystal-clear formal model
 - T: Global accurate time (starts from time 0)
 - For node A:


$$T_a = T * (1 + X_a) + Y_a$$

clock reading on A clock rate drift on A initial clock drift on A


X_a may change from time to time

Cristian's Algorithm


- Assumptions
 - Propagation delay for request and reply is the same
 - Clock rate drift is 0



- $T1, T2, T3, T4$ are accurate times
- $Ta2 - Ta1 = T4 - T1$ (since clock rate drift is 0)
- $Tb2 - Tb1 = T3 - T2$
- Propagation delay $d = ((T4 - T1) - (T3 - T2)) / 2$



- Key idea: $(T4 - T1)$ is accurate even if A has clock drift
- What is A's clock rate drift is not 0?
 - Consider a scenario where $d = 0$, $Tb2 - Tb1 = 1\text{ms}$, $Ta2 - Ta1 = 2$ minutes
- Thus this **implicit** assumption (not mentioned in textbook) is critical !



Network Time Protocol (NTP)

- Inaccuracy in Cristian's protocol comes from
 - Non-identical propagation delay for request and reply (quite common in wide-area network)
 - Clock rate drift not being 0 -- impact is usually small (why?)
- Total propagation delay is an upper bound on the difference in propagation delays for request and response
 - $d_1 + d_2 = d$ implies $|d_1 - d_2| \leq d$
- Synchronize with multiple servers and pick the one with smallest d

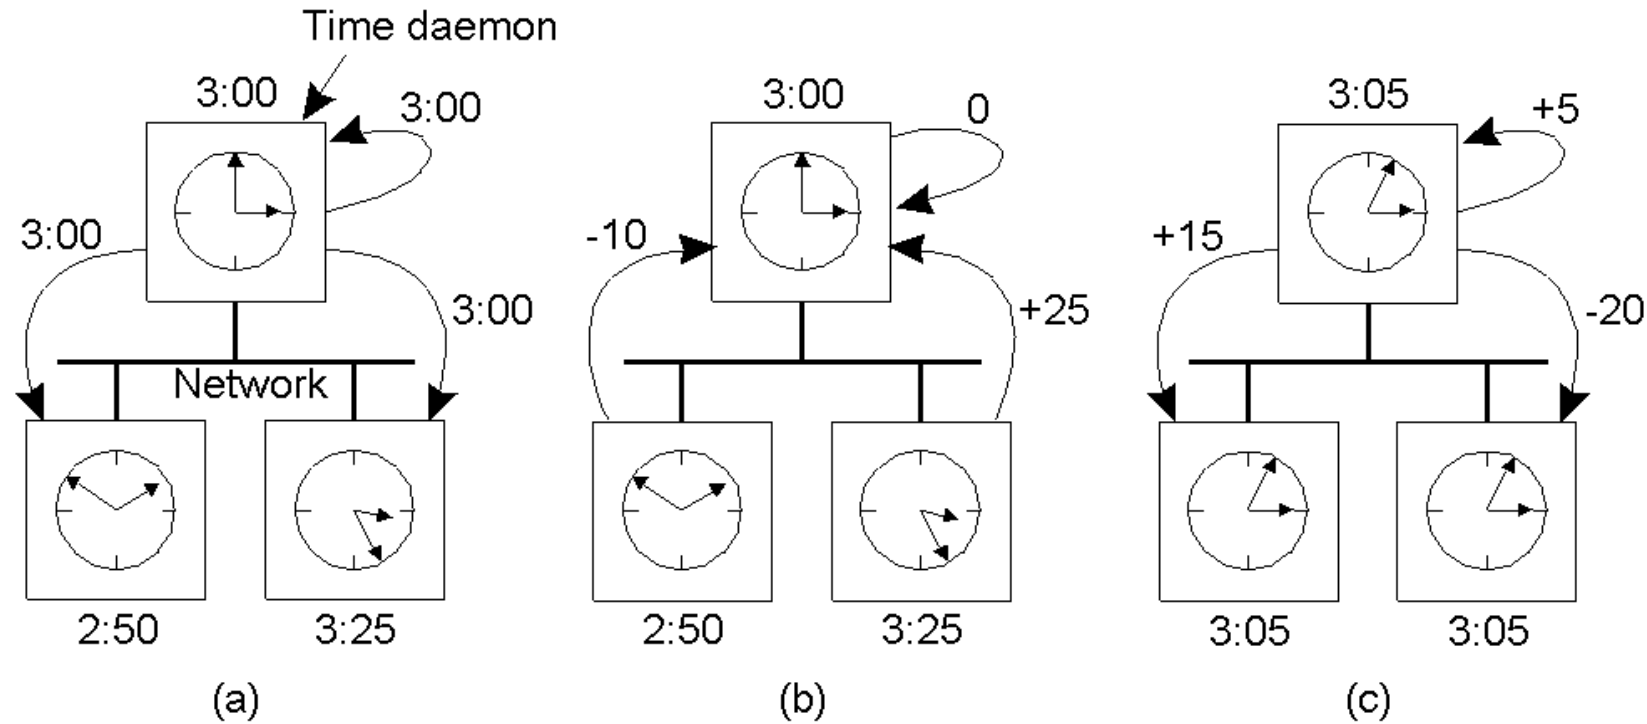
NTP Continued

- Does not make sense to synchronize with a inaccurate clock
- Stratum-1 servers:
 - Servers with UTC receiver or atomic clocks
- Machines synchronize with Stratum-1 servers become stratum-2
 - And so on...
- Current worldwide NTP accuracy
 - 1 to 50 milliseconds

The Berkeley Algorithm

- Assumptions
 - Network delay = 0
 - Usually for cases where no machine has the “accurate time”
- Idea: Hope that the clock drifts on different machines will cancel out

The Berkeley Algorithm: Example



The Berkeley Algorithm

- Quite straight-forward
- If you have a machine with “accurate time”
 - Should not take average
 - Just broadcast the time periodically

Problems with Physical Clocks

- Many protocols need to impose an ordering among events
 - If two players open the same treasure chest “almost” at the same time, who should get the treasure?
- Physical clocks:
 - Seems to completely solve the problem
 - But what about theory of relativity?
 - Even without theory of relativity – efficiency problems
- How accurate is sufficient?
 - Without out-of-band communication: Minimum message propagation delay
 - With out-of-band communication: distance/speed of light
 - In other words, some time it has to be “quite” accurate

Software “Clocks”

- Software “clocks” can incur much lower overhead than maintaining (sufficiently accurate) physical clocks
 - But does not allow comparison with external clocks
- Allows a protocol to infer ordering among events
- Goal of software “clocks”: Capture event ordering that are visible to users
 - Users do not have accurate physical clocks either
 - But what orderings are visible to users without physical clocks?


Assumptions

- Process can perform three kinds of atomic actions/events
 - Local computation
 - Send a single message to a single process
 - Receive a single message from a single process
 - No atomic broadcast
- Communication model
 - Point-to-point
 - Error-free, infinite buffer
 - Potentially out of order

Visible Ordering to Users

- $A < B$ (process order)
- $B < C$ (send-receive order)
- $A < C$ (transitivity)
- $A ? D$
- $B ? D$

PollEv.com/haifengyu229




“Happened-Before” Relation

- “Happened-before” relation captures the ordering that is visible to users when there is no physical clock
 - A partial order among events
 - Process-order, send-receive order, transitivity
- First introduced by Lamport – Considered to be the first fundamental result in distributed computing
- Goal of software “clock” is to capture the above “happened-before” relation


Software “Clock” 1: Logical Clocks

- Each event has a single integer as its logical clock value
 - Each process has a local counter C
 - Increment C at each “local computation” and “send” event
 - When sending a message, logical clock value V is attached to the message. At each “receive” event, $C = \max(C, V) + 1$



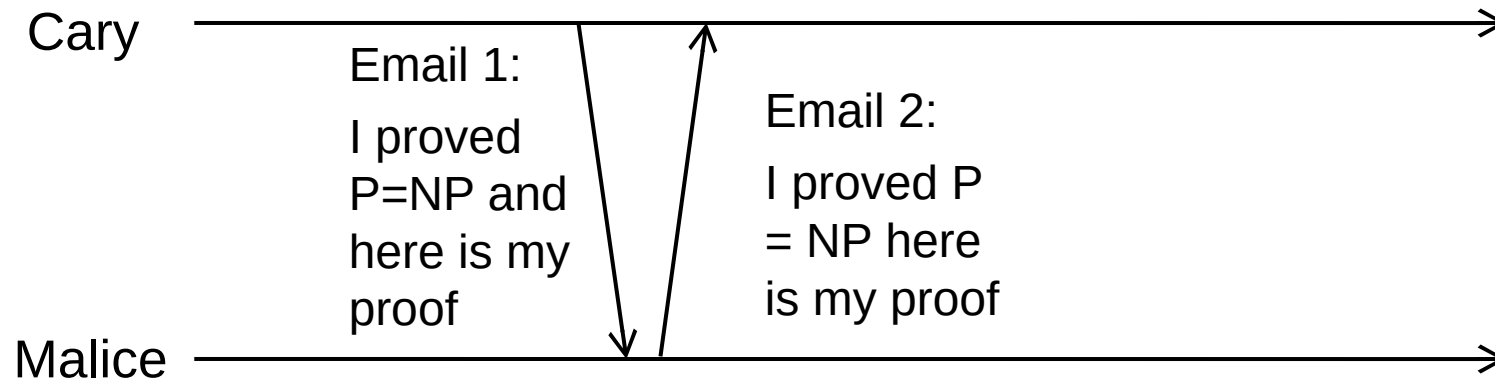
Logical Clock Properties

- Theorem:
 - Event s happens before $t \Rightarrow$ the logical clock value of s is smaller than the logical clock value of t .
- The reverse may not be true



Example Application for Logical Clock

- Total-ordered broadcast protocol is the “standard” example application for logical clock (as in textbook)
 - But it is complex and will take half a lecture to go through
 - (The protocol is covered in CS4231...)
- We will use a much simpler application example



- We want to decide who to give Turing award to
 - Both Cary and Malice try to claim credit
 - Malice says Cary copies her proof
 - Cary says Malice copies her proof


Software “Clock” 2: Vector Clocks

- Logical clock:
 - Event s happens before event t \Rightarrow the logical clock value of s is smaller than the logical clock value of t .
 - Vector clock:
 - Event s happens before event t \rightarrow the vector clock value of s is “smaller” than the vector clock value of t .
 - Each event has a vector of n integers as its vector clock value
 - $v1 = v2$ if all n fields same
 - $v1 \leq v2$ if every field in $v1$ is less than or equal to the corresponding field in $v2$
 - $v1 < v2$ if $v1 \leq v2$ and $v1 \neq v2$
- Relation “ $<$ ” here is not a total order

Vector Clock Protocol


- Each process i has a local vector C
- Increment $C[i]$ at each “local computation” and “send” event
- When sending a message, vector clock value V is attached to the message. At each “receive” event, $C = \text{pairwise-max}(C, V)$; $C[i]++$;

$C = (0,1,0)$, $V = (0,0,2)$
 $\text{pairwise-max}(C, V) = (0,1,2)$



Vector Clock Properties

- Event s happens before $t \Rightarrow$ vector clock value of $s <$ vector clock value of t : There must be a chain from s to t
- Event s happens before $t \Rightarrow$ vector clock value of $s <$ vector clock value of t
 - If s and t on same process, done
 - If s is on p and t is on q , let VS be s 's vector clock and VT be t 's
 - $VS < VT \Rightarrow VS[p] \leq VT[p] \Rightarrow$ Must be a sequence of message from p to q after s and before t




Example Application of Vector Clock

- Each Bitcoin node has some blocks (assume that Bitcoin runs on a fixed set of nodes)
 - Want all nodes to know all blocks
- Each block has a vector clock value


I have seen all blocks whose vector clock is smaller than $(2,3,2)$:

I can avoid linear search for existence testing




Mutual Exclusion: Token-based Approach

- Motivation is the same as in non-distributed systems
- Token-based approach
 - A single token system-wide
 - Whoever holds the token has the privilege
- Typical software architecture
 - A special node acts as “lock manager” or “token manager”



**Process 1
requests**



**Process 2
requests
lock while
process 1
is holding
the lock**

Mutual Exclusion: Token-based Approach

- Also possible to implement it as a token ring
- Pros:
 - Simplicity – lock manager approach is widely used
 - Fairness easy to ensure
 - No starvation
- Cons:
 - Cannot deal with failures nicely

Mutual Exclusion: Token-based Approach

- Dealing with client failures
 - Use leases – the lock is granted for only a finite amount of time
 - Automatically revoked after time out
 - Problem: A client is writing to a large file but cannot renew the lease due to lost network connection
- Dealing with server (lock manager) failures
 - Need to start a new server
 - Need to ensure the number of token is exactly 1

Mutual Exclusion: Voting

- We want to deal with server crashes
- Have n servers instead of 1
 - Each server has a vote
 - A server may crash and recover
- Assume clients do not crash
 - We can always use lease to deal with client crashes
- To acquire a lock, a client needs to get $(n/2+1)$ votes
 - Servers get the votes back after the client releases the lock
 - Called majority voting

Mutual Exclusion: Voting

- Theorem:
 - No two clients can hold the lock at the same time with majority voting (**even if some server crashes**)
- Can tolerate at most $n/2-1$ server failures
 - Usually quantified as availability
- Disadvantage:
 - Need n servers instead of one
 - Need to contact $n/2+1$ servers to get lock

Mutual Exclusion: Voting

- Majority voting is extremely flexible
 - How about assigning different number of votes to different servers?
 - A (5 votes), B(2 vote), C(2 vote), D(2 vote)
 - What is the impact on availability and performance?
- What is the best vote assignment mechanism?
 - Machines fail iid with probability of p
 - Theorem: Democracy is the best and monarchy is the worst when $p < 0.5$
 - Theorem: Democracy is the worst and monarchy is the best when $p > 0.5$

Mutual Exclusion: Quorum Systems

- A (5 votes), B(2 vote), C(2 vote), D(2 vote) =
{ {A, B}, {A, C}, {A, D},
{A, B, C}, {A, B, D}, {A, C, D}, {B, C, D},
{A, B, C, D} }
- A voting system can always be expressed as a quorum system
 - The reverse is not true
- But it happens that the best quorum systems corresponds to majority voting (under the earlier failure model)

History Readings (Non-compulsory)

- “The reliability of voting mechanisms”, IEEE Transactions on Computers, pages 1197-1208, Oct 1987.
- “How to assign votes in a distributed system”, Journal of ACM, October 1985.

Leader Election: Motivation

- We often need a coordinator in distributed systems
 - Leader, distinguished node/process
- We will only discuss leader election algorithms for failure-free contexts
 - Fault-tolerant leader election algorithms are related to fault-tolerance and will be covered later in the course

Leader Election on General Graph (n known)

- Complete graph
 - Each node send its id to all other nodes
 - Wait until you receive n ids
 - Biggest id wins
- Any connected graph
 - **Flood** your id to all other nodes (how?)
 - Wait until you receive n ids
 - Biggest id wins

Leader Election on General Graph (n unknown)


- Complete graph
 - n must be known
- Any connected graph
 - An auxiliary protocol to calculate the number of nodes
- The protocol is initiated by any node who wants to know n
 - The protocol actually establishes a spanning tree starting from the initiator

Spanning Tree Construction

- Remember: No centralized coordinator
 - Goal of the protocol: Each node knows its parent and children (i.e., a distributed tree)
- The root is initially marked
- Every marked node send “mark” msg to all its neighbors
- Upon receiving the first “mark” msg, an unmarked node becomes marked, and sends back “I am your child” msg
- For later “mark” msgs, send back “I am not your child” msg
- Tree construction completed if every node has received all replies from all its neighbours

Counting Nodes Using a Spanning Tree

- Disseminate “start count” msg down the tree
- Count will be aggregated up the tree



Summary

- Physical clock synchronization
- Software clocks and applications
- Mutual exclusion
- Leader election