

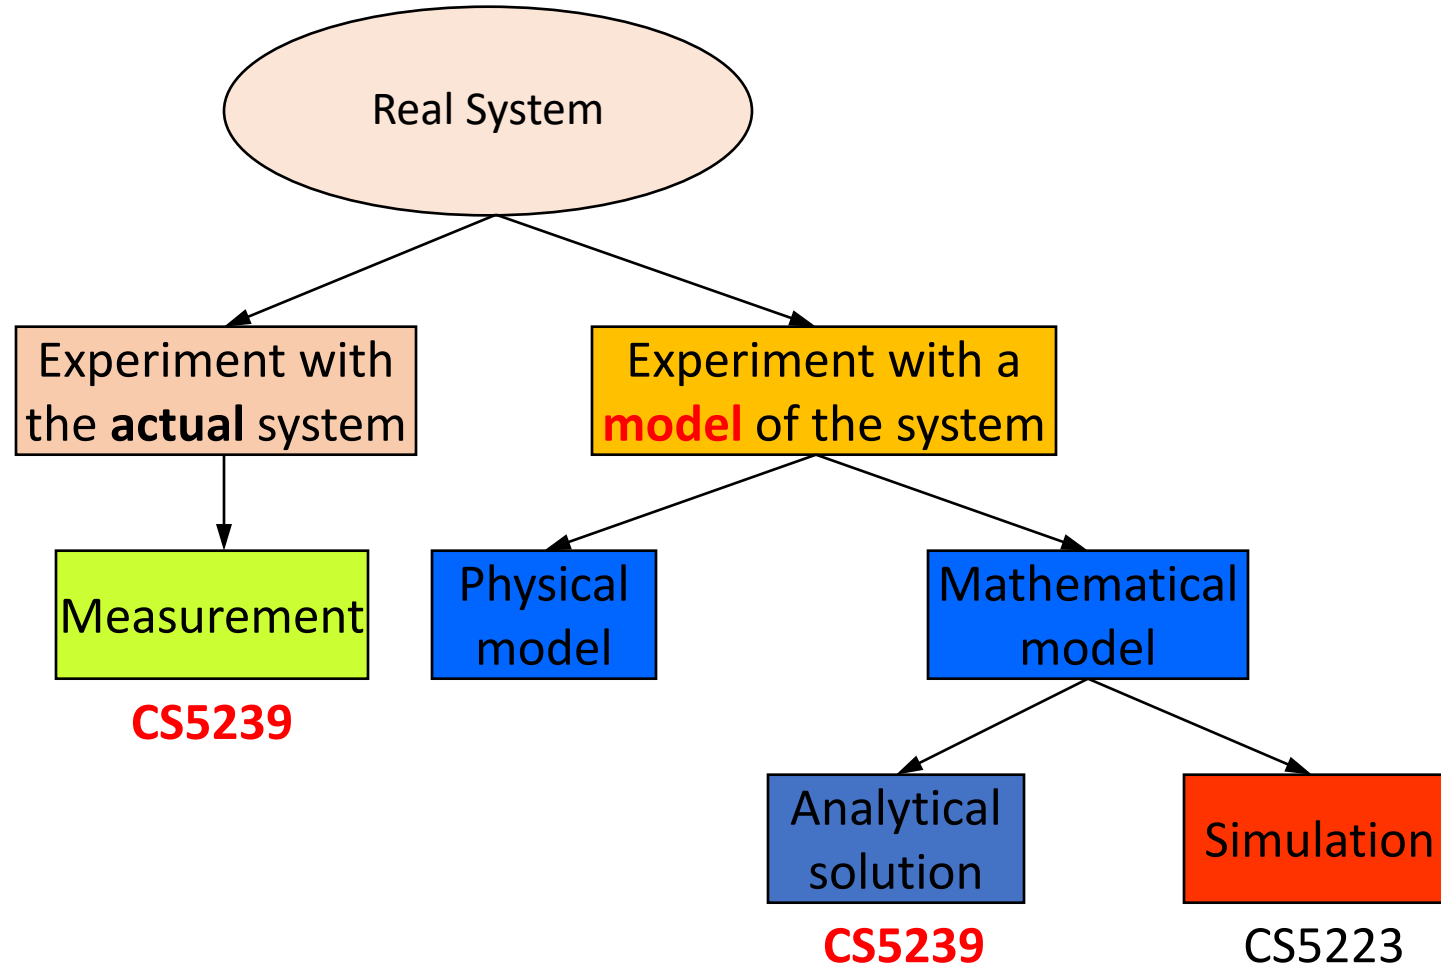
CS5239 Computer System Performance Analysis

Lecture 1 - Introduction


L01: Introduction

- What is performance?
- Performance scope and challenges
- Activities
- Latency
- Observability

Ways to Study a System



What is Performance?



- Shows how well a computer system performs a given job or activity
 - Latency (time)
 - Bandwidth (rate)
- Performance is important to achieve:
 - Time deadline
 - Cost (and efficiency) of computing resources
 - Energy cost

Performance of a computer system

- Performance of a computer system is multidimensional
 - complex component interactions
 - hard to predict how it will scale
 - ...
- Sharing of resource -> congestion/queueing



Performance Evaluation

- Performance evaluation is of interest to computer system **users**, administrators, developers, and designers
- Goal: obtain or provide **highest performance** at the **lowest cost**

performance-cost ratio

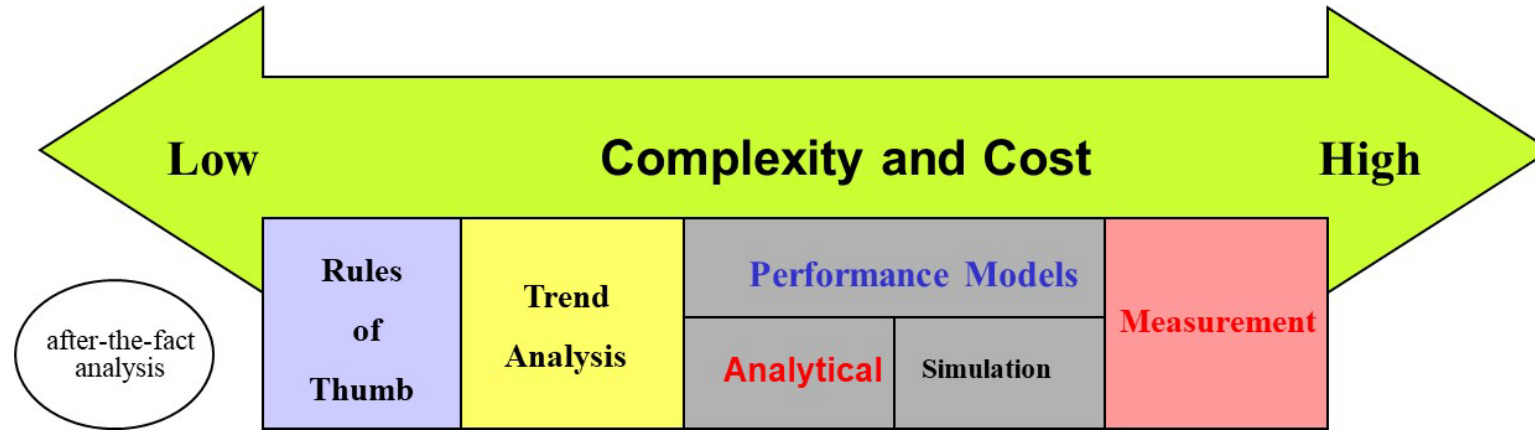
Trade-offs

- The good/fast/cheap “pick two” trade-off



- Many IT projects choose on-time and inexpensive, leaving performance to be fixed later


Performance Evaluation: How



- Measurements of actual systems
- Modeling
 - Simulation using software models
 - Analytical modeling using techniques such as queueing analysis

Performance Evaluation Techniques

- Analytical modelling
 - Mathematical representation of the system
 - Model assumptions may not be met in practice - lower accuracy
 - Ease of changing parameters
 - Key insights – validation of simulations / measurements
- Simulation
 - Program written to model important features of the system
 - Can be easily modified to study impact of changes
 - Cost – writing and running the program
 - Impossible to model every small detail – sometimes low accuracy



Models are abstraction of the system

Performance Evaluation Techniques

- Measurement
 - No simplifying assumptions
 - Highest credibility
 - Only works if systems / applications exist
 - Information only on specific system being measured
 - Harder to change system parameters in a real system
 - May be intrusive and perturb operation of systems
 - Difficult and time consuming
 - Need for software tools

Goals of Performance Evaluation

- Set expectations → service level
 - Provide information for users
- System tuning & performance debugging
 - Find parameters that produce best overall performance
 - Identify bottlenecks
 - Determine impact of a new feature
- Compare alternative designs
 - Which configurations are best under which conditions
 - Which program/algorithm is faster
- ...



Applications of Performance Evaluation

- Selection/configuration of computer systems:
 - System upgrade
 - System design: design of applications and equipment
- Analysis of existing systems: **capacity planning** predicts when future load levels will saturate the system and determines cost-effective way of delaying system saturation

[Global Internet Traffic](#)

Scalability

- The performance of the system under increasing load
 - Linear scalability until the **knee point**
 - A component reaches 100% utilization: **saturation point**
 - Queueing is frequent




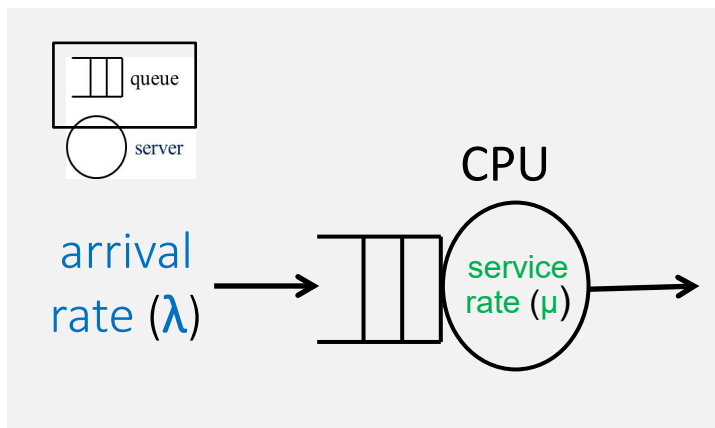
Known and Unknowns

US Secretary of Defence Donald Rumsfeld, 2002

- There are **known knowns**; there are things we know we know
 - Example: you know that you should check the CPU utilization and you know utilization is 10%
- We also know there are **known unknowns**; that is to say we know there are some things we do not know.
 - Example: you know that you could be checking what is making the CPU busy but have yet to do so
- But there are also **unknown unknowns** – there are things we do not know we don't know.
 - Example: you may not know what device interrupts are making the CPU busy and thus you are not checking them
- Performance issues can originate from anywhere, including areas of the system that you know nothing about and therefore not checking (unknown unknowns)

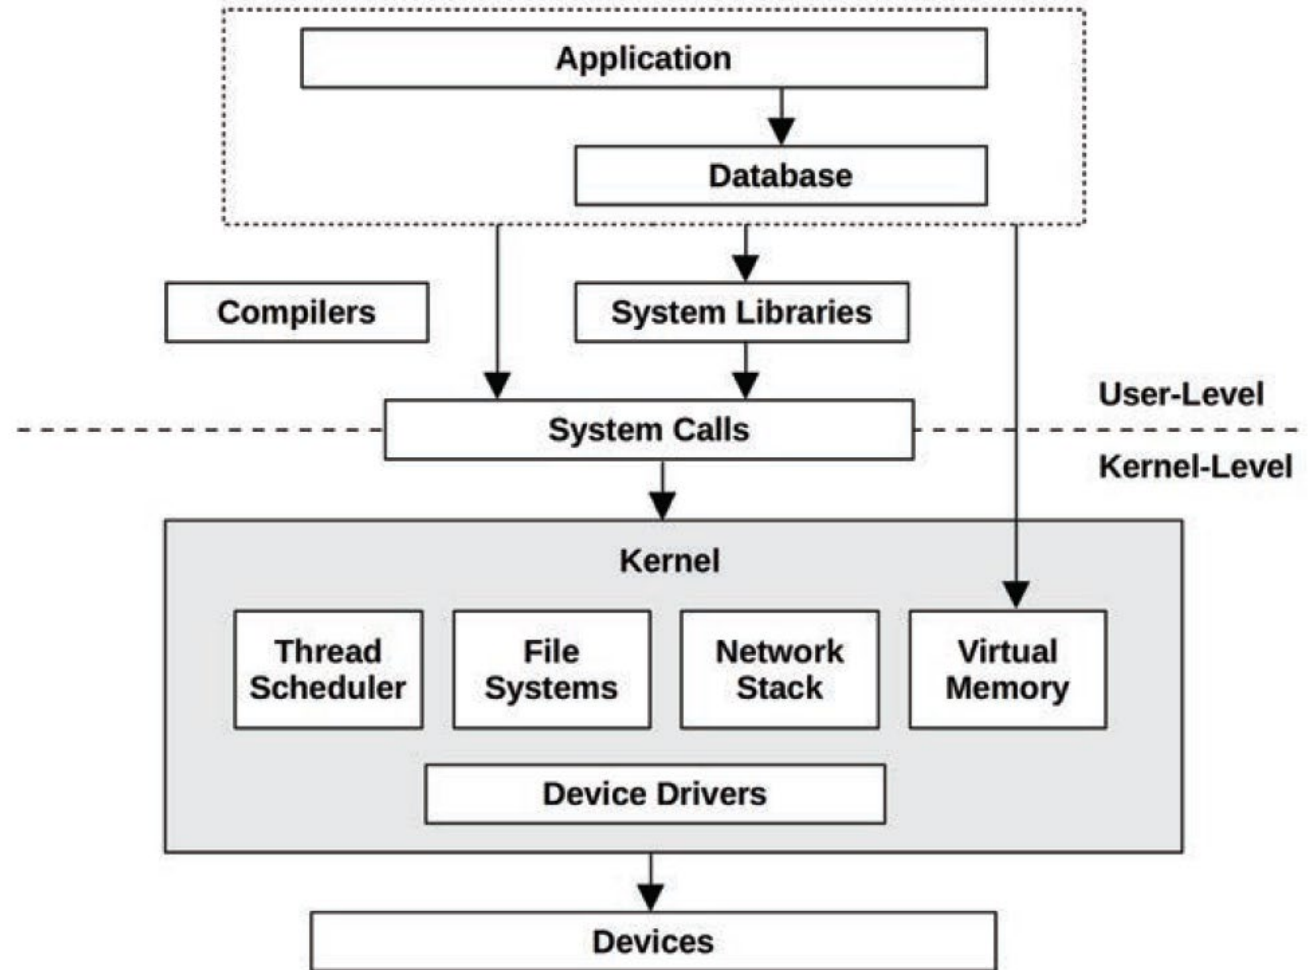
Performance Questions

1. What is the **average time** it takes a job to complete service? [**latency**]
2. What is the **throughput** of the system (number of jobs completed per unit time)? [**bandwidth**]
3. If **arrival rate is doubled** ($\lambda \rightarrow 2\lambda$), do we do nothing or how much should the **service rate** (μ) increase?
4. If we need more server (CPU) capacity, what are our options?
 - a. Buy a **new server** with the needed capacity
 - b. Buy a **few smaller servers** that adds up to the required capacity
 - c. How do we organize these servers?
 - i. One queue for all servers
 - ii. One queue for each server
 - iii. Does it matter?



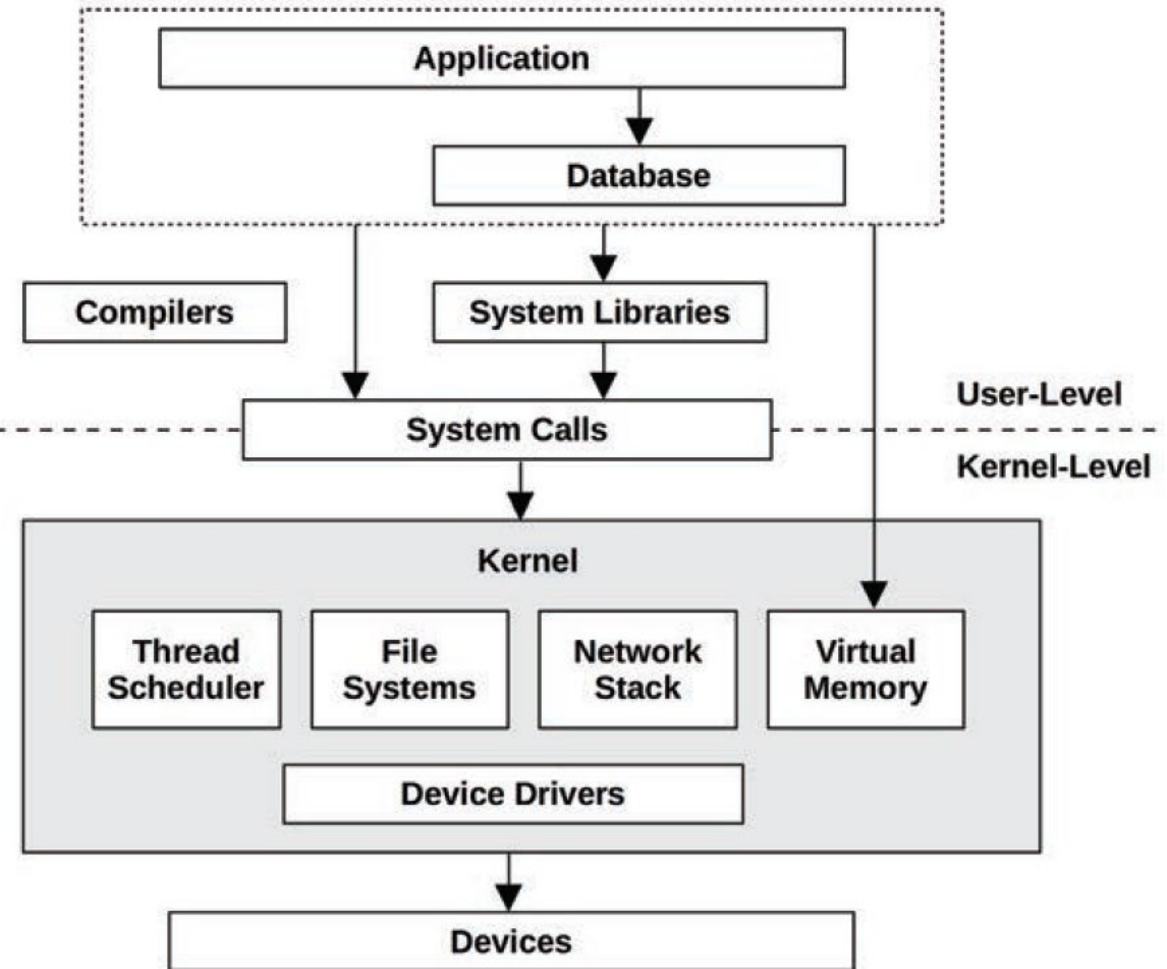
Systems Performance

- Study the performance of an entire computer system
 - All major software and hardware components
 - Anything in the data path, from storage devices to application software
 - For distributed systems, multiple servers and applications



Systems Performance

- Study the performance of the full system stack
- **Full stack** → the entire software stack from the application down to metal (the hardware)
 - Includes system libraries, the kernel, and the hardware
 - Includes the compiler



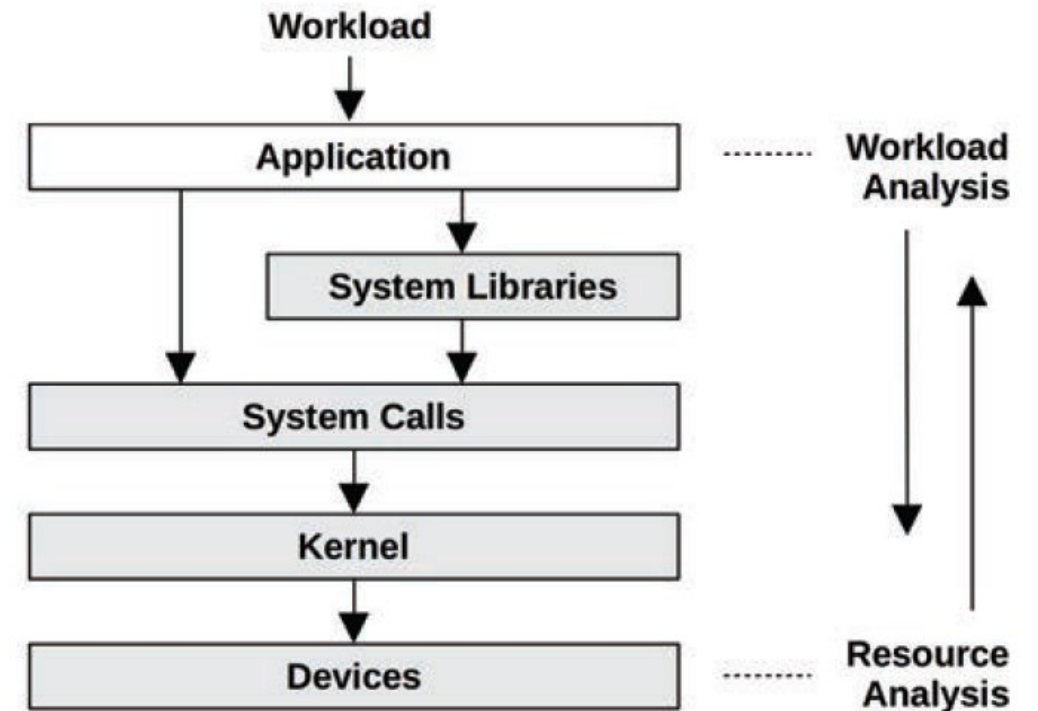
Performance Activities

- Should **follow the life cycle** of a software project from conception through development to production deployment
- Set objectives and create a **performance model** as a first step in product development
 - Don't defer performance engineering work to a later time, after a problem arises
- **Capacity planning** – study the resource footprint of development software to see how well the design can meet the target needs
- Cloud applications performance
 - *Canary testing*
 - *Blue-green development*

Perspectives

- The resource analysis perspective
→ system administrators
- The workload analysis perspective
→ application developers

Operating
System
Software
Stack



Challenges

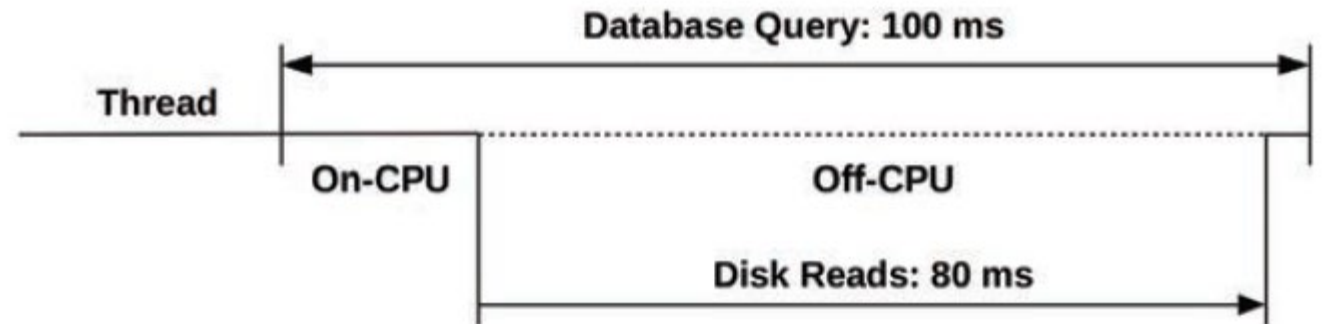
- Subjectivity
 - “The average disk I/O response time is 1 ms.” → is it good or bad?
- Complexity
 - Lack of an obvious starting point for analysis
 - Start with a hypothesis
 - Cascading failure – one failed component causes performance issues in others
 - Bottlenecks – complex and unexpectedly related to each other
- Multiple causes: issues do not have a single root cause
- Multiple performance issues
 - Quantify the magnitude of issues

Metrics

- Metrics are needed to **quantify** the performance of a system
- **Latency** - measure of time spent waiting
 - Broadly, it means the time for any operation to complete
 - Time taken by an application request, a database query, a file system operation, etc.
 - Example: the time needed for a website to load completely, from link click to screen paint
 - Allows maximum speedup to be estimated
 - Ambiguous term without qualifying terms
 - In networking, latency can mean the time for a connection to be established but not the data transfer time [connection latency]; or it can mean the total duration of a connection, including the data transfer [request latency]

Latency - Example

- A database query that takes 100 ms [**latency**] during which it spends 80 ms blocked waiting for disk reads
- The maximum performance improvement is obtained by eliminating disk reads (e.g., by caching) from 100 ms to 20 ms
 - 5x faster [**speedup**]



- What if you are looking at I/O operations per second instead?

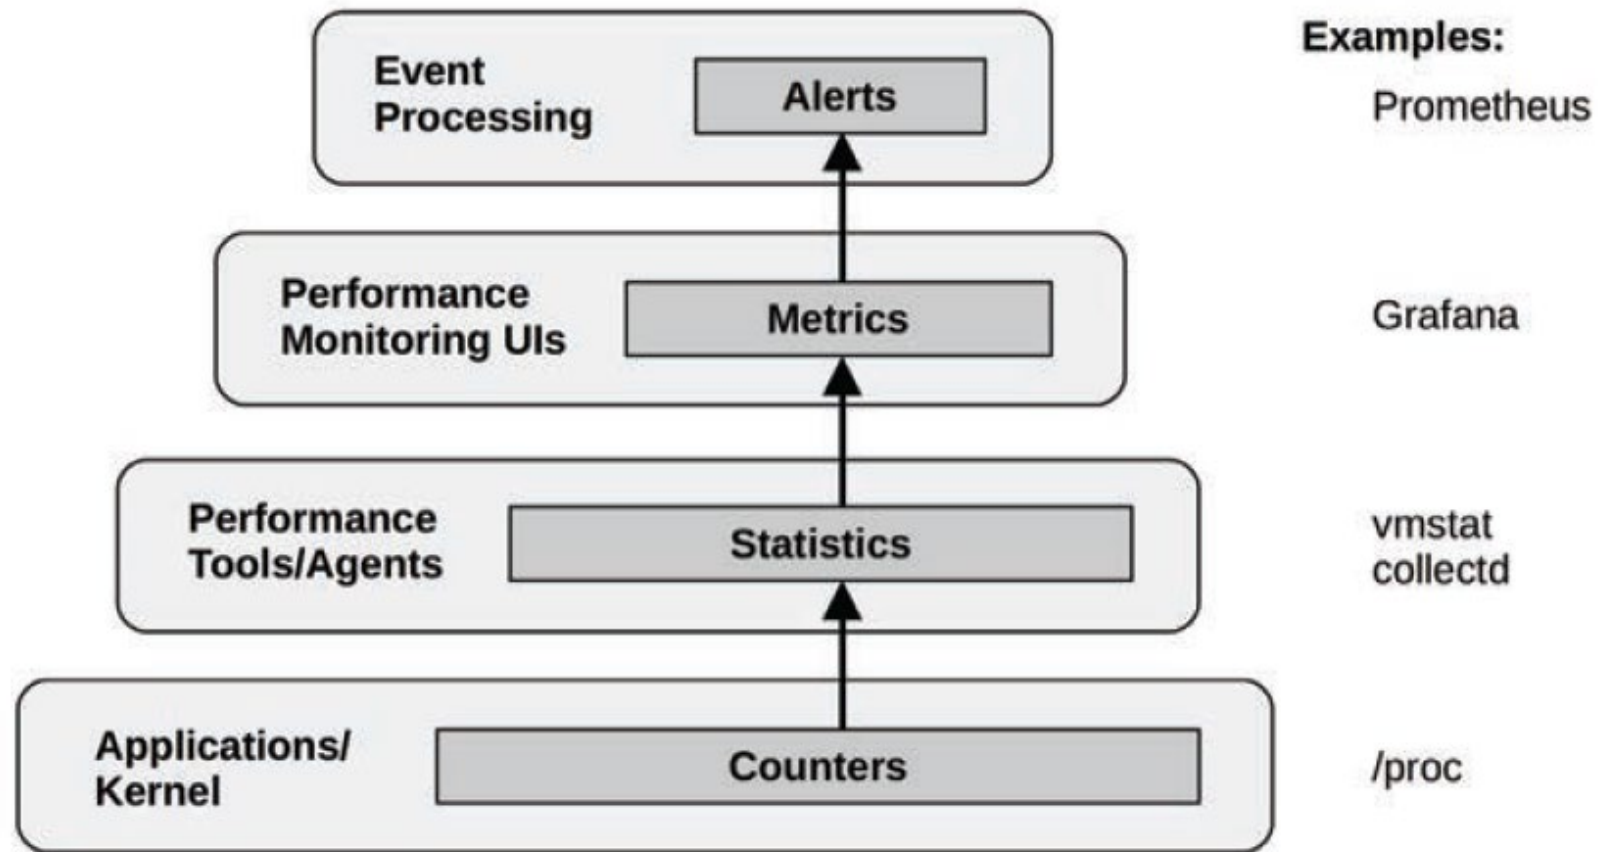
Observability

- Understanding a system through observation using different observability tools
- Observability tools use
 - Counters
 - Profiling
 - Tracing
- For production environments, observability tools may perturb production workloads through resource contention

Counters, Statistics, Metrics, and Alerts (1)

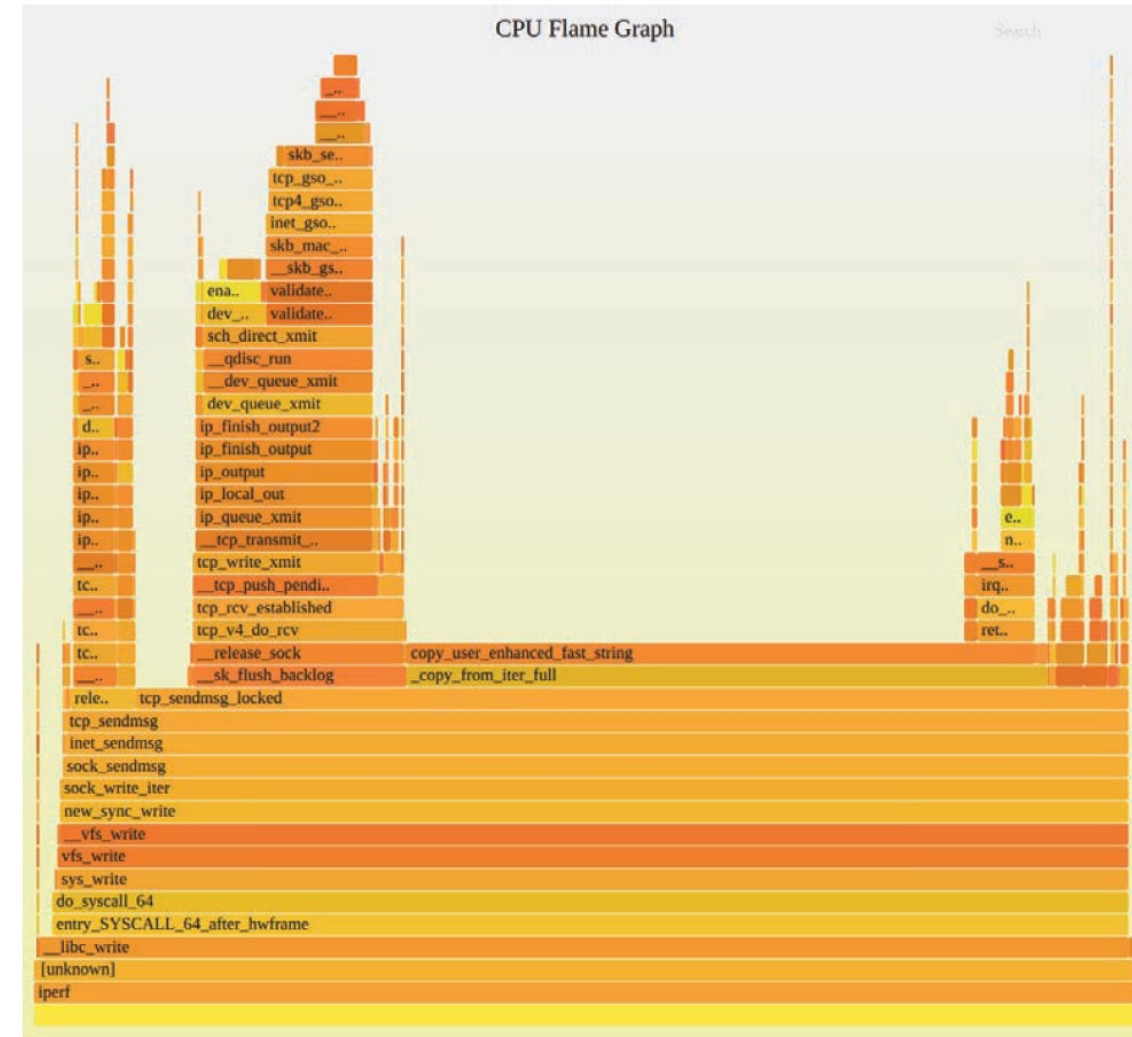
- Applications and the kernel typically provide data on their state and activity:
 - Example: operation counts byte counts, latency measurements, resource utilization, and error rates
 - Implemented as integer variables called **counters**
 - Hard-coded in the software
 - Some are cumulative, and they always increment
- Cumulative counters can be read at different times by performance tools for calculating **statistics**
 - Examples: the rate of change over time, the average, percentiles, etc.
 - **Metric** is a statistic that has been selected to evaluate or monitor a target
 - Monitoring software can also support creating custom **alerts** from these metrics

Performance Instrumentation Terminology



Profiling

- Refers to the use of tools that perform sampling: taking a subset (a sample) of measurements to paint a coarse picture of the target
 - CPUs are a common profiling target
 - Take timed-interval samples of the on-CPU code paths
- Visualizing the profiles is important
 - Example: CPU flame graphs




Tracing

- Event-based recording: event data is captured and saved for later analysis or consumed on-the-fly for custom summaries
 - Special-purpose tracing tools for system calls (Linux strace(1)) and network packets (e.g., Linux tcpdump(8))
 - General-purpose tracing tools that can analyze the execution of all software and hardware events (e.g., Linux Ftrace, BCC, and bpftrace).
- Static instrumentation: hard-coded software instrumentation points added to the source code
- Dynamic instrumentation creates instrumentation points after the software is running, by modifying in-memory instructions to insert instrumentation routines

Steps in Performance Evaluation Study


1. State the goals of the study and define the system **boundaries**
2. List system **services** and possible **outcomes**
3. Select performance **metrics**
4. List system and workload **parameters**
5. Select **factors** and their values
6. Select evaluation **techniques**
7. Select the **workload**
8. Design the **experiments**
9. Analyze and interpret the **data**
10. Present the **results**. Start over, if necessary



Example: Performance Evaluation Study

1. System definition:

- compares performance of applications using remote pipes (caller is not blocked) with remote procedure calls (calling program is blocked)



Example (2)

2. Service:

- two types of channel calls: remote procedure call and remote pipe

3. Metrics:

- resources are local computer (client), remote computer (server), and network link
- performance metrics:
 - a. Elapsed time per call
 - b. Maximum call rate per unit of time, or equivalently, the time required to complete a block of n successive calls
 - c. Local CPU time per call
 - d. Remote CPU time per call
 - e. Number of bytes sent on the link per call

Example (3)

4. Parameters

- i. workload parameters (resource demand)
 - a. Time between successive calls
 - b. Number and sizes of the call parameters
 - c. Number and sizes of the results
 - d. Types of channel
 - e. Other loads on the local and remote CPUs
 - f. Other loads on the network
- ii. System parameters (resource capacity)
 - a. Speed of the local CPU
 - b. Speed of the remote CPU
 - c. Speed of the network
 - d. Operating system overhead for interfacing with the channels
 - e. Operating system overhead for interfacing with the network
 - f. Reliability of the network affecting the number of retransmissions required

Example (4)

5. Key Factors:

- a. Types of channel: 2 types – compares remote pipes and remote procedure calls
- b. Speed of the network: 2 remote host locations - short distance (in the campus) and long distance (across the country)
- c. Sizes of call parameters to be transferred: small and large
- d. Number n of consecutive calls: 11 values of n – 1, 2, 4, 8, 16, 32,..., 512, 1024

Factors are selected based on resource availability and objectives of evaluation

Example (5)

6. Evaluation Technique: Since both types of channels are available, measurements is used for evaluation. Analytical modeling is used to justify the consistency of measured values for different parameters.
7. Workload: a synthetic program generating the specified types of channel requests, monitor resources consumed and log the measured results.
8. Experimental Design: a full factorial experimental design consists of 88 (23 x 11) experiments
9. Data Analysis: Analysis of Variance to quantify the effects of the first three factors and regression to quantify the effects of the number n of successive calls
10. Data Presentation: final results plotted as a function of the block size n

Summary

- Motivation for performance analysis
- Tradeoffs
- Metrics
- Observability
- Steps in performance analysis

References

- [Gregg] – chapter 1
- [Jain] – chapters 2 & 3