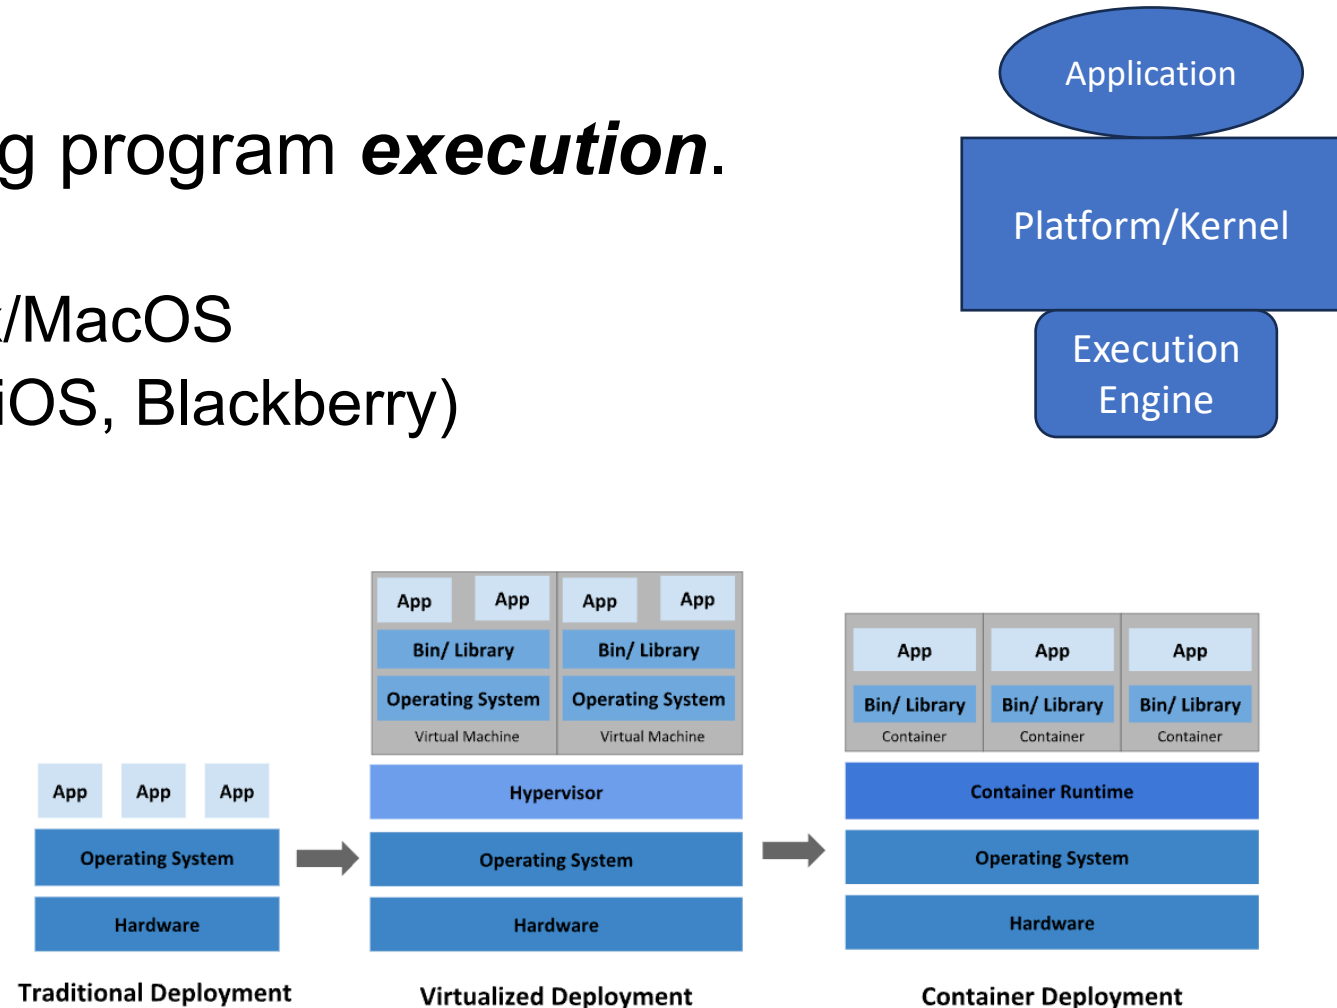


CS5231: Systems Security

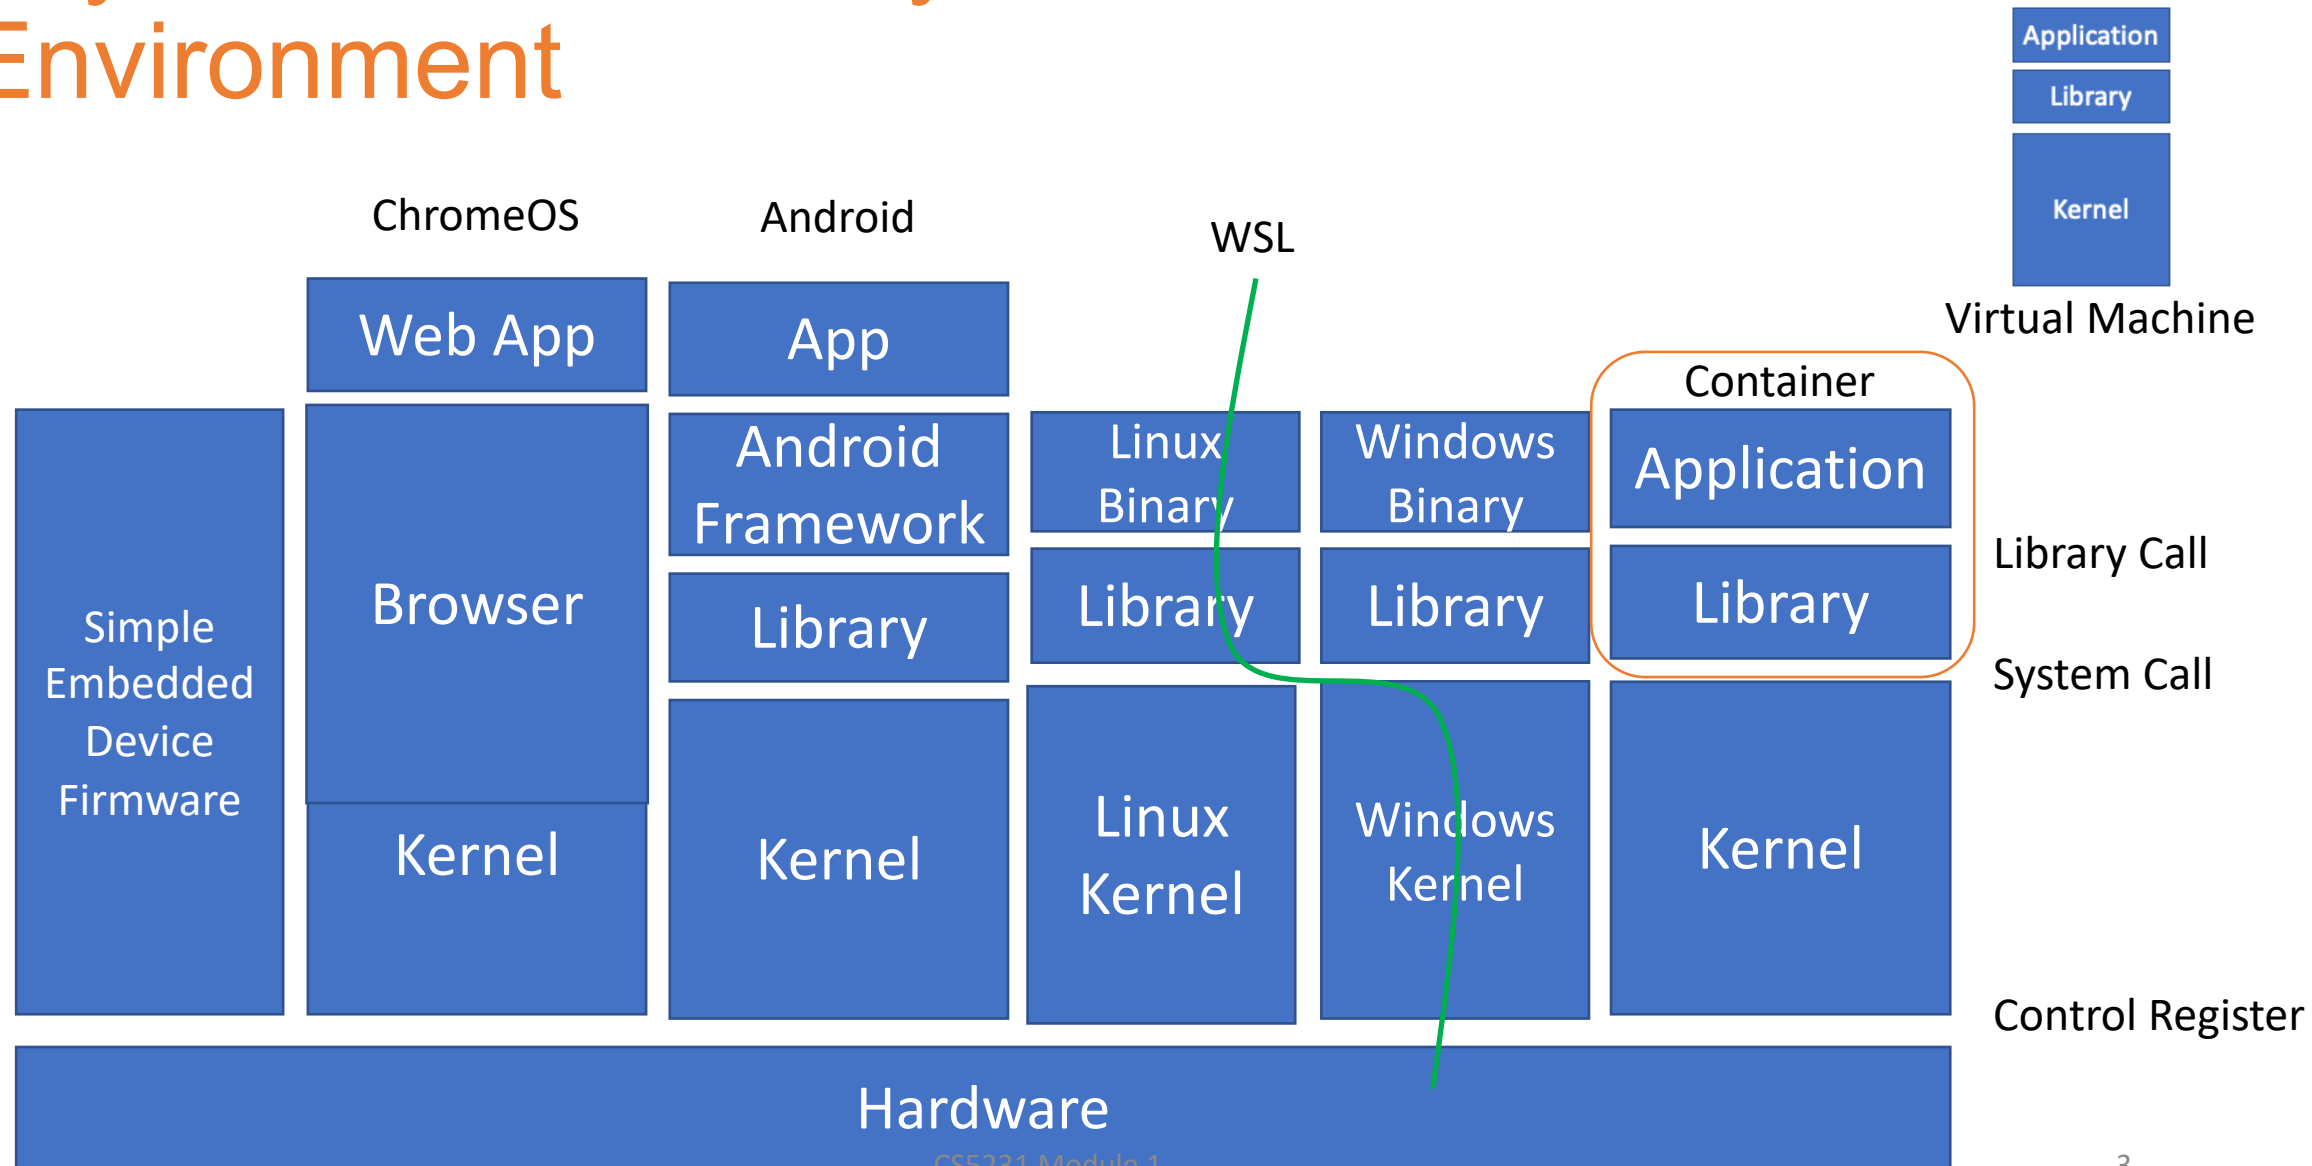
Module 1: Memory Safety Vulnerabilities (Part 1) Basis

What is a System?

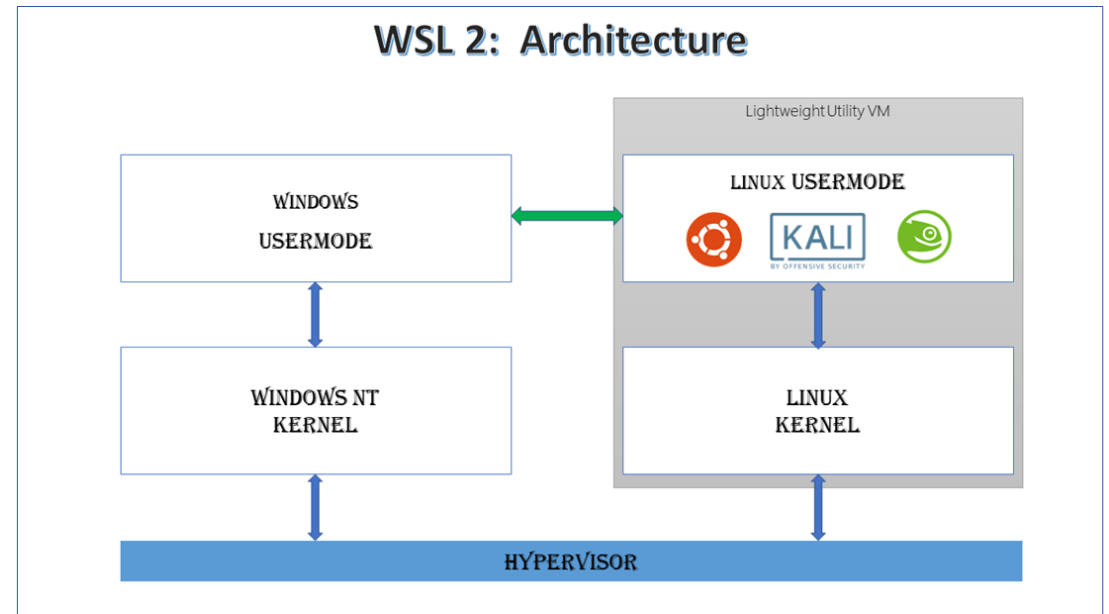
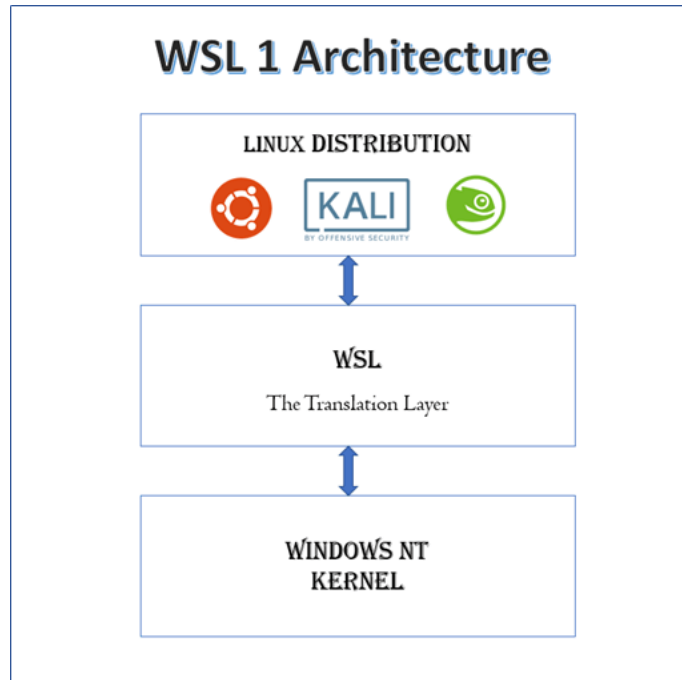
- A **platform** supporting program **execution**.
 - Embedded system
 - DOS, Windows/Linux/MacOS
 - Mobile OS (Android, iOS, Blackberry)
 - Browser
 - Virtual Machine
 - Container
 - Cloud
 - ...



Layers and Flexibility of Execution Environment



Windows Subsystem for Linux (WSL)



Discussion

- How to run a Linux GUI program on WSL?
 - Kernel
 - System call
 - Library
 - X Window server

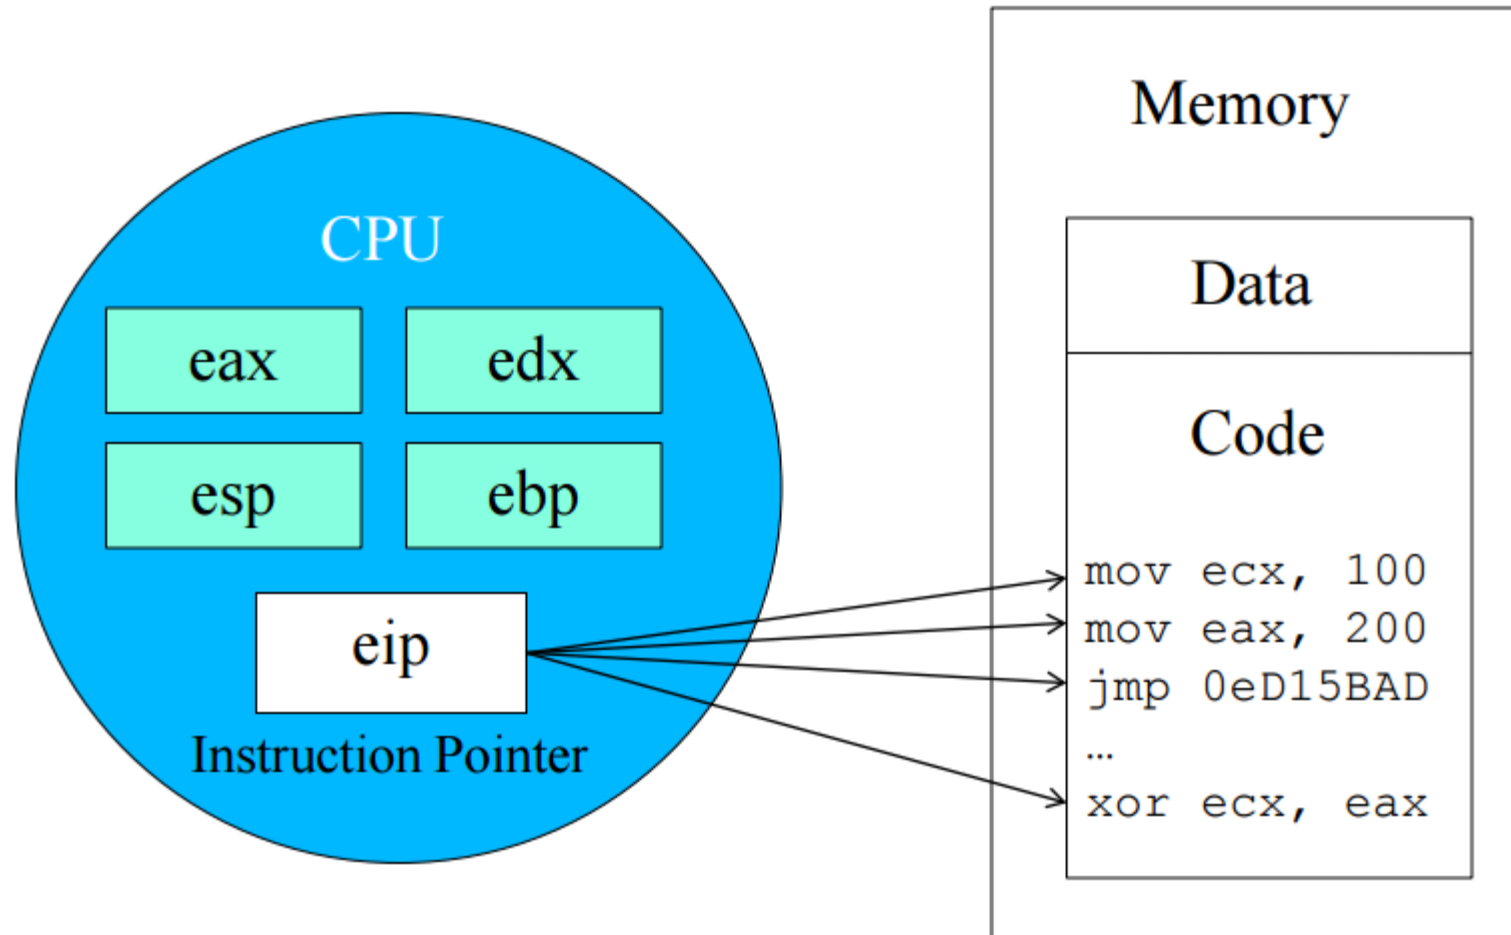
Semester Plan

- Module 1: Memory error and defense
- Module 2: System auditing and provenance
- Module 3: Kernel security

- Midterm: Week 8 (October 11)
- Quiz: take-home online quiz

Execution on CPU

Basics: The x86 Machine Model



Basics: The x86 Machine Model


- Both code and data are represented as numbers

- Code

- `lea ecx, [esp+4]` represented as `0x8d 0x4c 0x24 0x04`


- Data

- On Intel CPUs, least significant bytes is put at lower addresses
 - It is called little endian
 - For example, `0x01020304`



Basics: The x86 Machine Model

- Registers, Instructions, Stack, EIP
- Addressing modes, offset addresses
 - `mov 0x12[ebp], ecx`
- Stack grows down, other memory accesses move up.



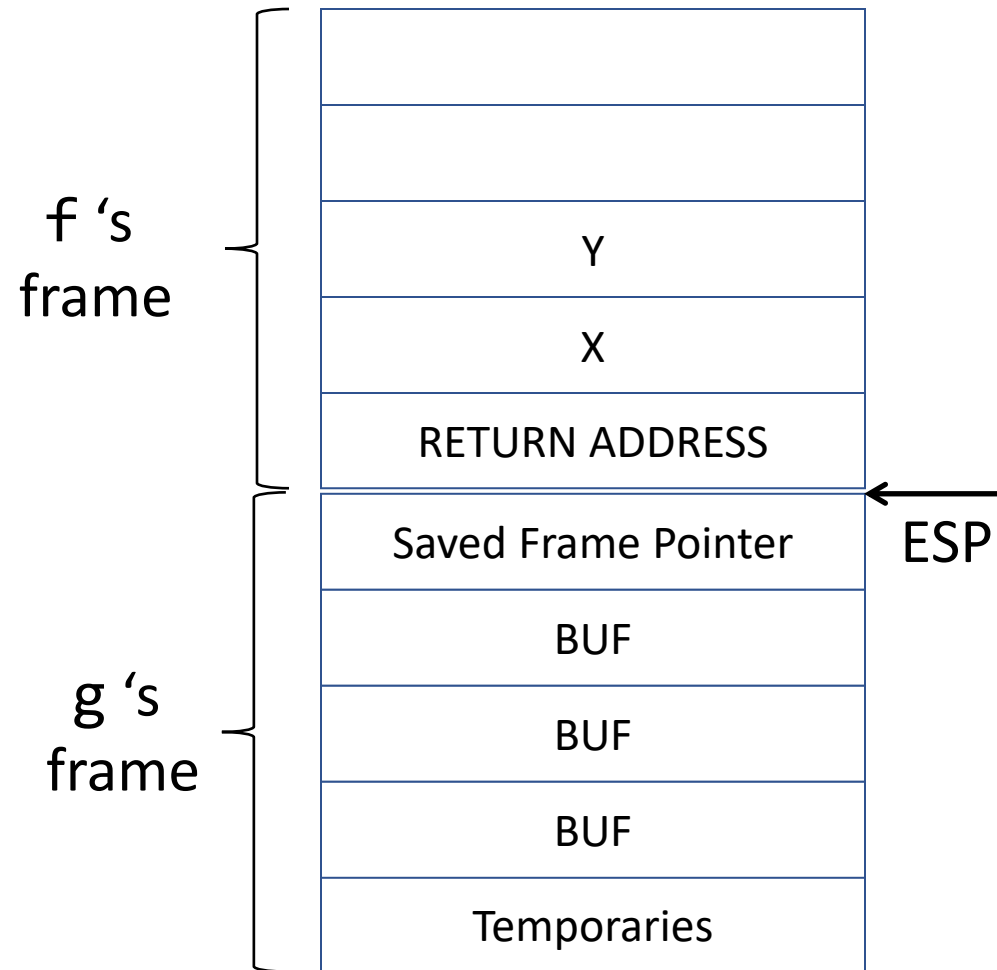
Stack Frames

```
int f() {  
...  
    g (x, y);  
}  
  
int g(int x, int y) {  
    char buf[50];  
    scanf("%s", buf);  
}
```

```
.g  
push ebp  
  
...  
call scanf  
...  
pop ebp  
ret
```



What address will it return to?



See It Using gdb

- Compile program and run gdb
 - `gcc -o sample -g sample.c`
 - `gdb sample`
- gdb commands
 - Set break point: `break <functionname>`
 - Check register values: `info registers`
 - Check variables: `print <variablename>`
 - Inspect memory: `x/b <variable_or_address>`

Spatial Memory Errors: Buffer Overflows


Buffer Overflows

```
int f() {  
    ...  
    g (x, y);  
}  
  
int g(int x, int y) {  
    char buf[50];  
    scanf("%s", buf);  
}
```

```
.g  
push ebp  
  
...  
call scanf  
...  
pop ebp  
ret
```



What address will it return to?

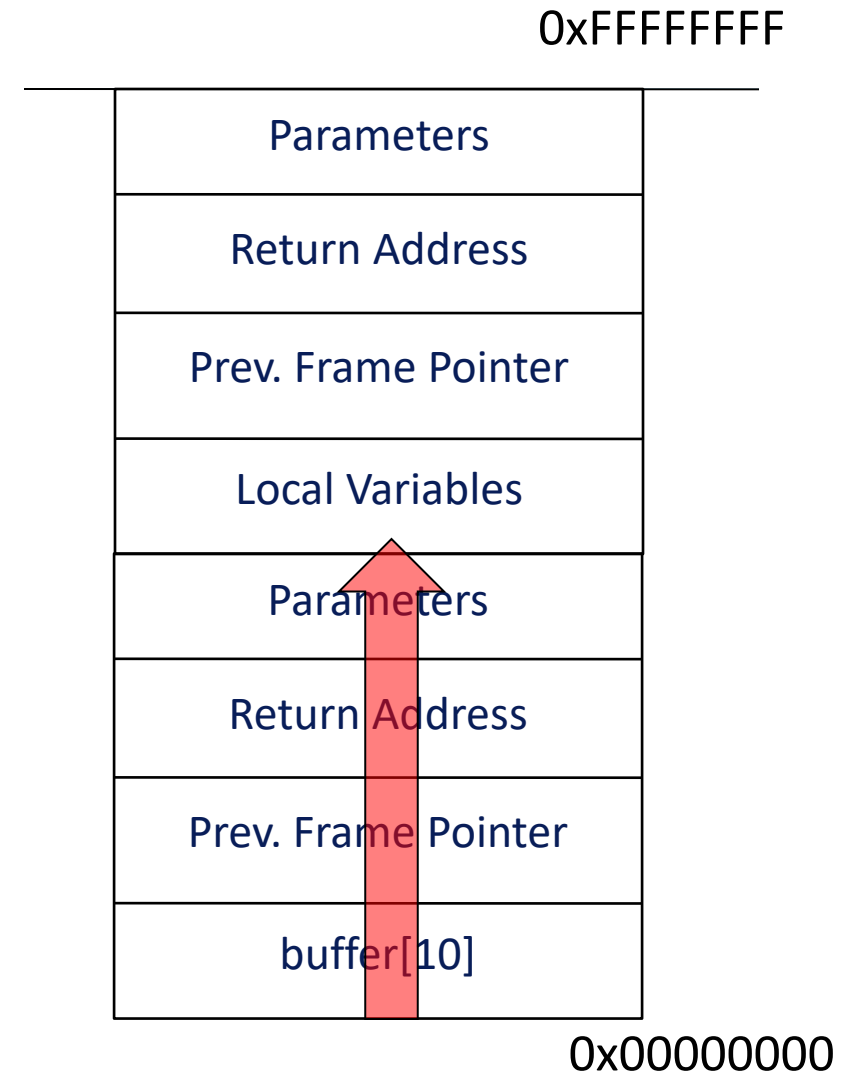


Buffer Overflow

```
void
sample_function(void)
{
    → char buffer[10];
    gets(buffer);
    return;
}


main()
{
    sample_function();
    printf("Loc 1\n");

    sample_function();
    printf("Loc 2\n");
}
```



Malicious Code Injection

- Remember executable code also represented as bytes
- Attackers can include code in the input
 - Called shell code
- They can arrange the return address to point to the injected code




Can They Know the Exact Address of Injected Code?

- Attackers can analyze the vulnerable program using a debugger to find out the address of target stack frame.
- Directly jump into system libraries
 - E.g., `system()` will execute a command
 - Where is its arguments?
- Sometimes, attacks only know the possible range of the code they injected

NOP Sled

- Instruction NOP, No Operation.
 - Tell CPU to do nothing and fetch the next instruction
- Including a large block of NOP instructions in the injected code as *landing area*
- Execution will reach shell code as long as return address pointing to somewhere in the NOP sled



Shell Code Example

```
int main(int argc,  
        char*argv[])  
{  
    char *sh;  
    char *args[2];  
  
    sh =  
    "/bin/bash";  
    args[0] = sh;  
    args[1] = NULL;  
    execve(sh,  
    args, NULL);  
}
```

- Shell Code

```
90 90 eb 1a 5e 31 c0 88 46 07  
8d 1e 89 5e 08 89 46 0c b0  
0b 89 f3 8d 4e 08 8d 56 0c  
cd 80 e8 e1 ff ff ff 2f 62  
69 6e 2f 73 68 20 20 20 20  
20 20
```

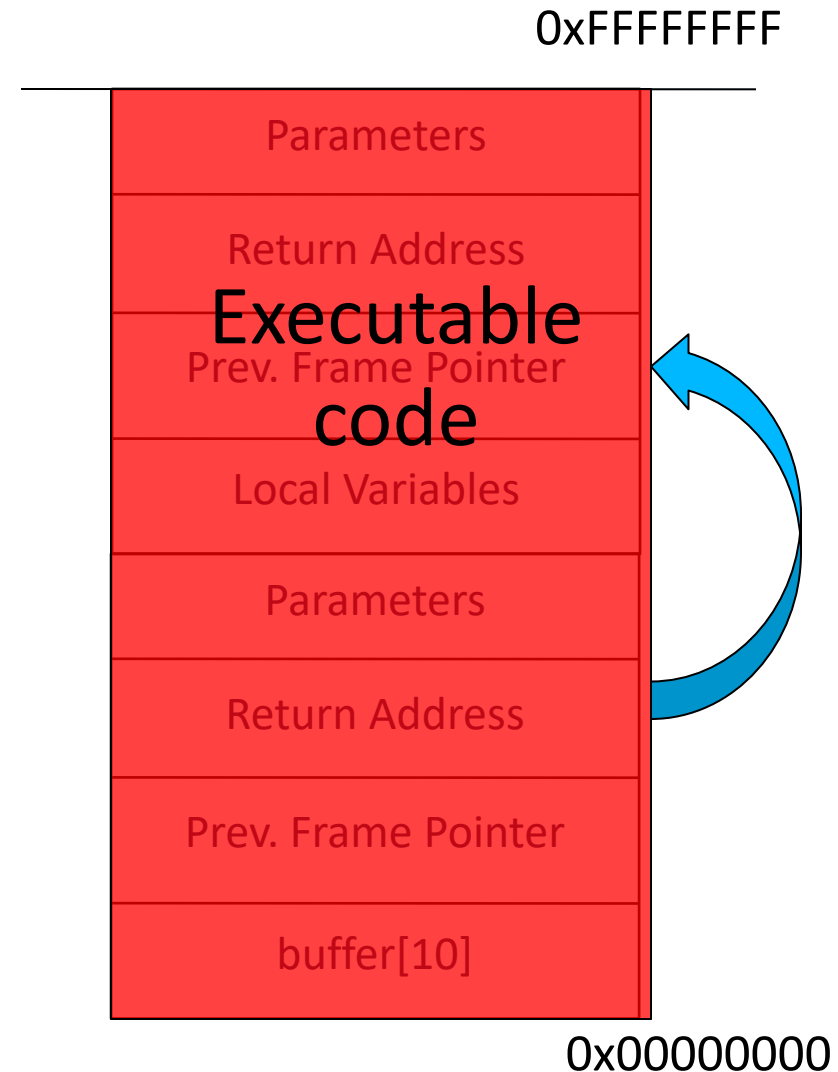
Targets of Buffer Overflow Attack

- Memories that control program execution
 - Return address
 - Function pointer
 - Virtual function table
- Important variables in program
 - Variable storing user id
 - Current balance of bank application
 - ...

Buffer Overflow Defense

Requirements of BO

- Existence of vulnerability
- Overwriting important data
- Known location of injected code
- Executable code in input



Essential Steps of Buffer Overflow

- Three essential steps
 - Attack Code Injection
 - How did you inject the attack code?
 - Control Flow Hijacking
 - How did you guess the address of attack code?
 - Attack Code Execution
 - How did you write the shellcode?

Buffer Overflow Defense (1)

- Existence of vulnerability
- Overwriting important data
- Known location of injected code
- Executable code in input

Safe Language and Coding

- Choose a safe programming language
 - Strong notion of variable types, such as Java
- Safe coding techniques
 - Pay attention to loops
 - Explicitly specify size of destination buffer
- Use safe libraries

Some Unsafe C Lib Functions

- `strcpy(char *dest, const char *src)`
- `strcat(char *dest, const char *src)`
- `gets(char *s)`
- `sprintf(const char *format, ...)`


- **Safe versions:**
- `strncpy, strncat, fgets, snprintf`

Code Checking Tools

- Tools checking for vulnerabilities using static source code analysis
 - ITS4 (It is the Software, Stupid --- Security Scanner)
 - RATS (Rough Auditing Tool for Security)
 - Flawfinder

Security Extension -- Libsafe

- Idea: Making unsafe functions safe!
 - Intercepts calls to `strcpy (dest, src)`
 - Validates sufficient space in current stack frame:
 $|\text{frame-pointer} - \text{dest}| > \text{strlen}(\text{src})$
 - If so, does `strcpy`.
Otherwise, terminates application.



Buffer Overflow Defense (2)

- Existence of vulnerability
- **Overwriting important data**
- Known location of injected code
- Executable code in input

MemGuard

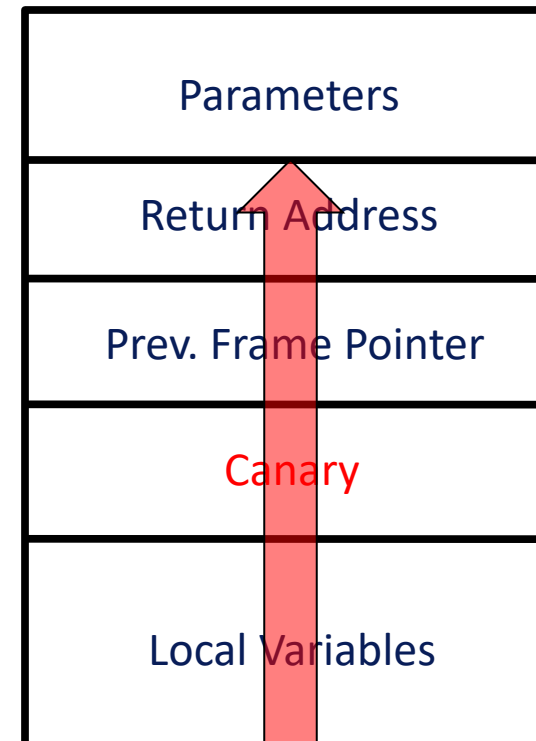
- Main idea: prevent return address from changing via mark it as **read-only**.
 - Extend **VM** model to protect user pages (such as return address in stack).
 - Ex: GCC's `function_prologue` and `function_epilogue`
- Flaw:
 - *Performance penalties*
 - Loading VM hardware is a privileged operation, and so the application process must trap to kernel mode to protect a word.

StackGuard

- Main idea: the technique used to smash the stack currently always involve **sequential memory writing**.
 - If the return address in stack was destroyed, the content before the return address must be destroyed, too.
 - Keep a “canary word” **before** return address and check this word before function returns
- Simple Demo
 - <http://nsfsecurity.pr.erau.edu/bom/StackGuard.html>

StackGuard

- Put a value below saved frame pointer upon entering the function, called canary
- Check canary value before function exit
- Changed canary value indicates an overflow
- Turned on by default in current gcc, try it out!



Flaw in static canary

- What if the attacker can easily guess the canary value?
- Workaround?
 - **randomize** canary-word

Buffer Overflow Defense (3)

- Existence of vulnerability
- Overwriting important data
- **Known location of injected code**
- Executable code in input

Why Randomization?

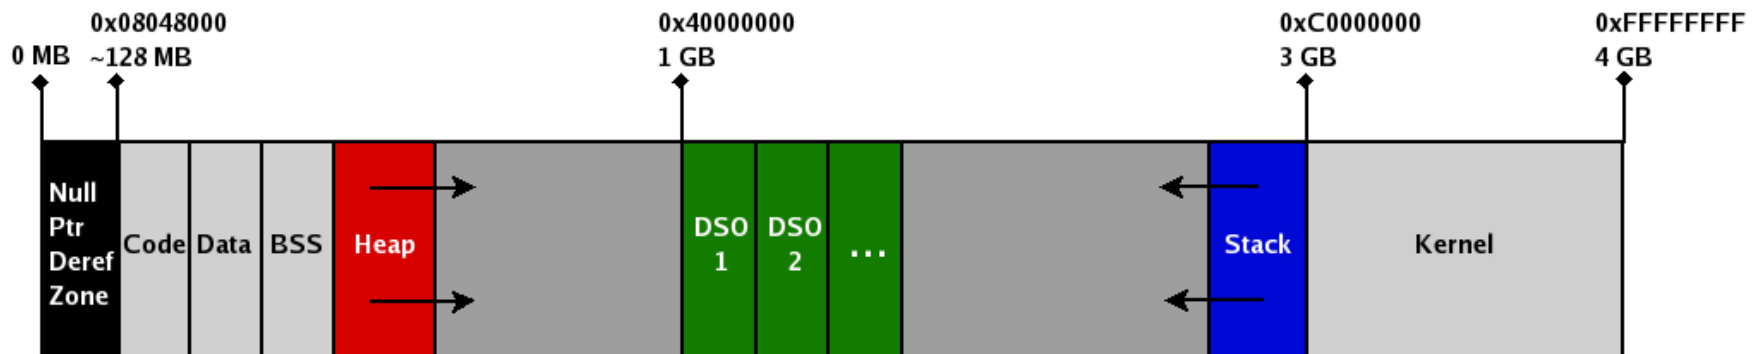
- Buffer overflow and **return-into-libc** exploits need to know the (virtual) address to which pass control
 - Address of attack code in the buffer
 - Address of a standard kernel library routine
- Same address is used on many machines
 - Slammer infected 75,000 MS-SQL servers using same code on every machine
- Idea: introduce **artificial diversity**
 - Make stack addresses, addresses of library routines, etc. unpredictable and different from machine to machine

Address Space Layout Randomization (ASLR)

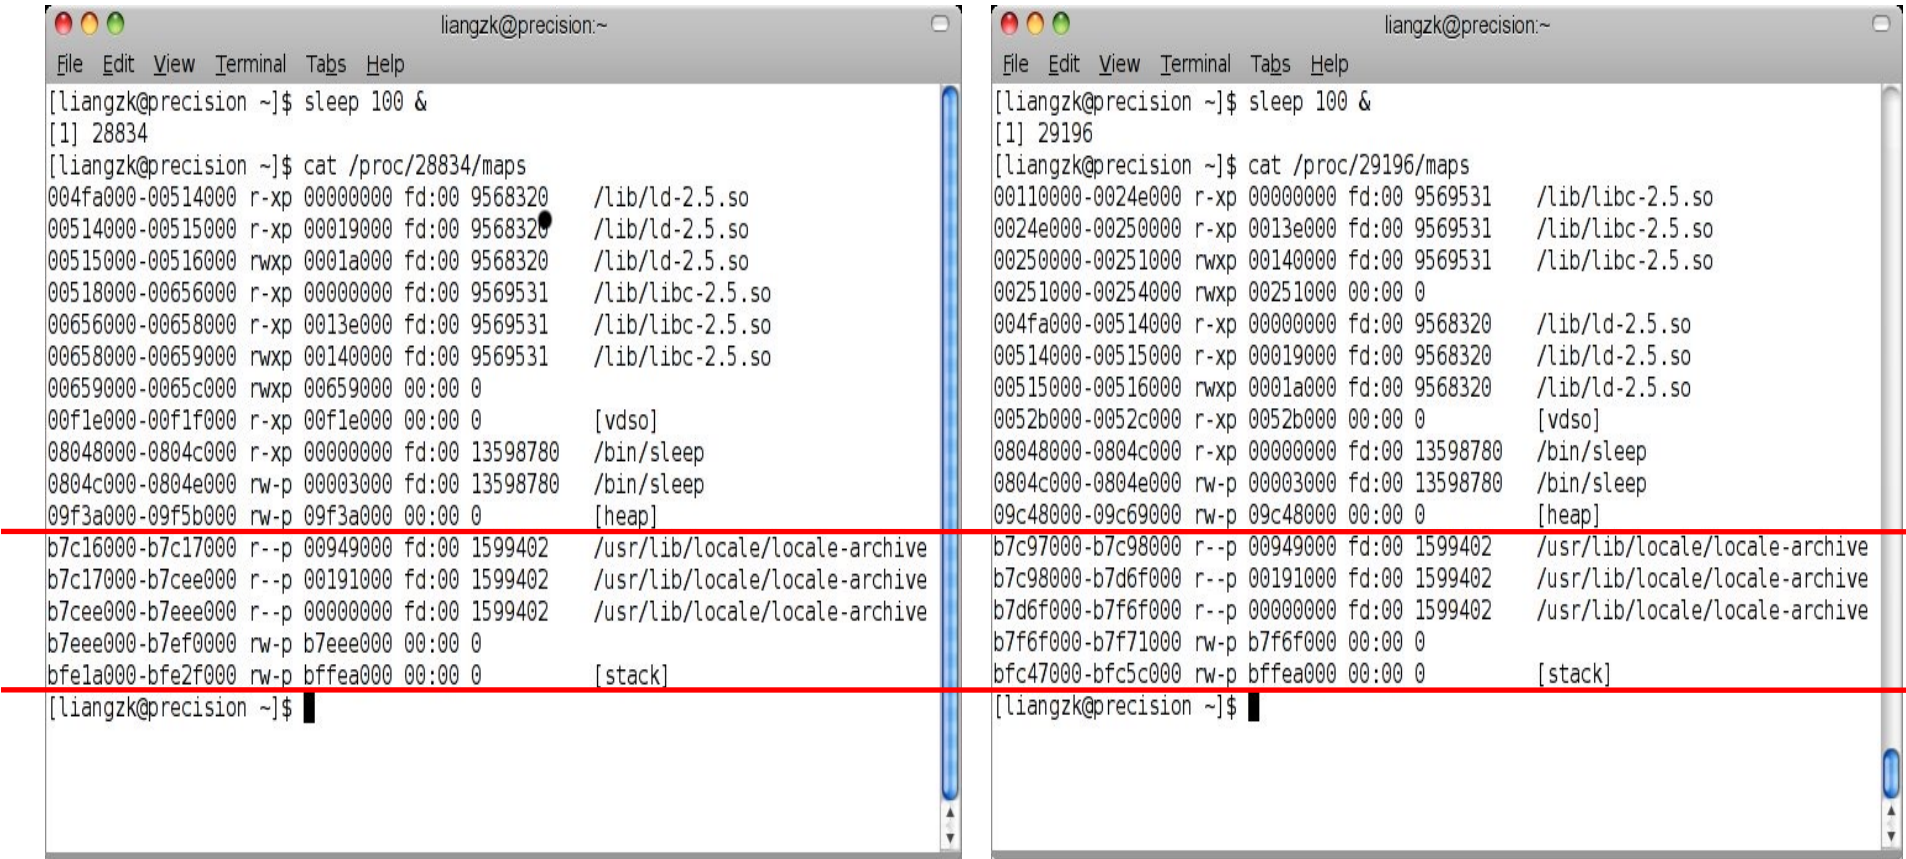
- ASLR renders exploits *which depend on predetermined memory addresses* useless by randomizing the layout of the virtual memory address space.
 - Base addresses of stack, heap, and code segment
- Randomization can be done at compile- or link-time, or by rewriting existing binaries
- Several implementations available
 - E.g., PaX ASLR

An Example: Linux VM System

- Each process has its own 32-bit address space
- Regions are page aligned
- Code & data regions fixed
- Multiple instances, same memory layout
- Traditional locations of heap, user stack, mmap
 - Randomization moves these three regions



ASLR in Linux



The image displays two terminal windows side-by-side, both titled 'liangzk@precision:~'. Each window shows the output of the command 'cat /proc/[PID]/maps' after running 'sleep 100 &'. The left window shows the memory map for PID 28834, and the right window shows the memory map for PID 29196. Both maps show randomized addresses for shared libraries like /lib/ld-2.5.so and /lib/libc-2.5.so, as well as the heap and stack. Red horizontal lines are drawn across the maps to highlight specific memory regions.

```
liangzk@precision:~  
File Edit View Terminal Tabs Help  
[liangzk@precision ~]$ sleep 100 &  
[1] 28834  
[liangzk@precision ~]$ cat /proc/28834/maps  
004fa000-00514000 r-xp 00000000 fd:00 9568320 /lib/ld-2.5.so  
00514000-00515000 r-xp 00019000 fd:00 9568320 /lib/ld-2.5.so  
00515000-00516000 rwxp 0001a000 fd:00 9568320 /lib/ld-2.5.so  
00518000-00656000 r-xp 00000000 fd:00 9569531 /lib/libc-2.5.so  
00656000-00658000 r-xp 0013e000 fd:00 9569531 /lib/libc-2.5.so  
00658000-00659000 rwxp 00140000 fd:00 9569531 /lib/libc-2.5.so  
00659000-0065c000 rwxp 00659000 00:00 0  
00f1e000-00f1f000 r-xp 00f1e000 00:00 0 [vdso]  
08048000-0804c000 r-xp 00000000 fd:00 13598780 /bin/sleep  
0804c000-0804e000 rw-p 00003000 fd:00 13598780 /bin/sleep  
09f3a000-09f5b000 rw-p 09f3a000 00:00 0 [heap]  
b7c16000-b7c17000 r--p 00949000 fd:00 1599402 /usr/lib/locale/locale-archive  
b7c17000-b7cee000 r--p 00191000 fd:00 1599402 /usr/lib/locale/locale-archive  
b7cee000-b7eee000 r--p 00000000 fd:00 1599402 /usr/lib/locale/locale-archive  
b7eee000-b7ef0000 rw-p b7eee000 00:00 0  
bfe1a000-bfe2f000 rw-p bffea000 00:00 0 [stack]  
[liangzk@precision ~]$
```

```
liangzk@precision:~  
File Edit View Terminal Tabs Help  
[liangzk@precision ~]$ sleep 100 &  
[1] 29196  
[liangzk@precision ~]$ cat /proc/29196/maps  
00110000-0024e000 r-xp 00000000 fd:00 9569531 /lib/libc-2.5.so  
0024e000-00250000 r-xp 0013e000 fd:00 9569531 /lib/libc-2.5.so  
00250000-00251000 rwxp 00140000 fd:00 9569531 /lib/libc-2.5.so  
00251000-00254000 rwxp 00251000 00:00 0  
004fa000-00514000 r-xp 00000000 fd:00 9568320 /lib/ld-2.5.so  
00514000-00515000 r-xp 00019000 fd:00 9568320 /lib/ld-2.5.so  
00515000-00516000 rwxp 0001a000 fd:00 9568320 /lib/ld-2.5.so  
0052b000-0052c000 r-xp 0052b000 00:00 0 [vdso]  
08048000-0804c000 r-xp 00000000 fd:00 13598780 /bin/sleep  
0804c000-0804e000 rw-p 00003000 fd:00 13598780 /bin/sleep  
09c48000-09c69000 rw-p 09c48000 00:00 0 [heap]  
b7c97000-b7c98000 r--p 00949000 fd:00 1599402 /usr/lib/locale/locale-archive  
b7c98000-b7d6f000 r--p 00191000 fd:00 1599402 /usr/lib/locale/locale-archive  
b7d6f000-b7f6f000 r--p 00000000 fd:00 1599402 /usr/lib/locale/locale-archive  
b7f6f000-b7f71000 rw-p b7f6f000 00:00 0  
bfc47000-bfc5c000 rw-p bffea000 00:00 0 [stack]  
[liangzk@precision ~]$
```

Buffer Overflow Defense (4)

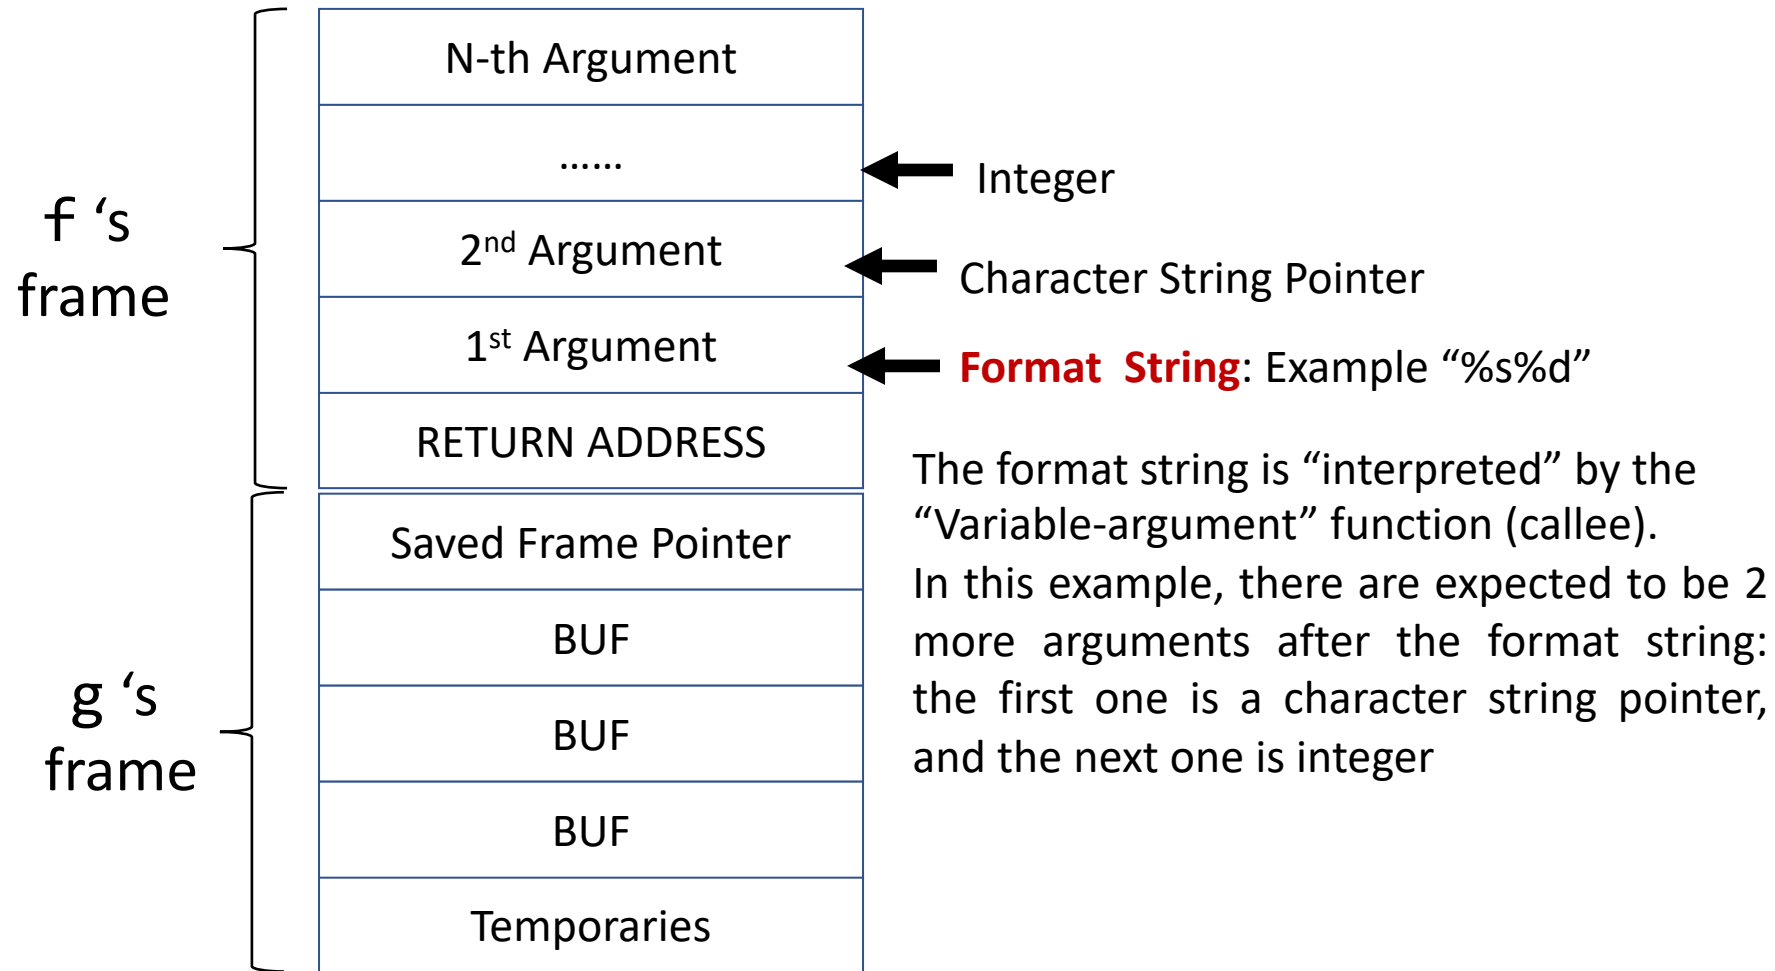
- Existence of vulnerability
- Overwriting important data
- Known location of injected code
- **Executable code in input**

Non-Executable Stack

- Using CPU's memory management unit to mark stack as non-executable
- Vulnerable program crashes if it jump to the stack for execution
- But there are legitimate reasons to put code on stack
 - Self-modifying code, e.g., Skype
 - Linux signal handlers

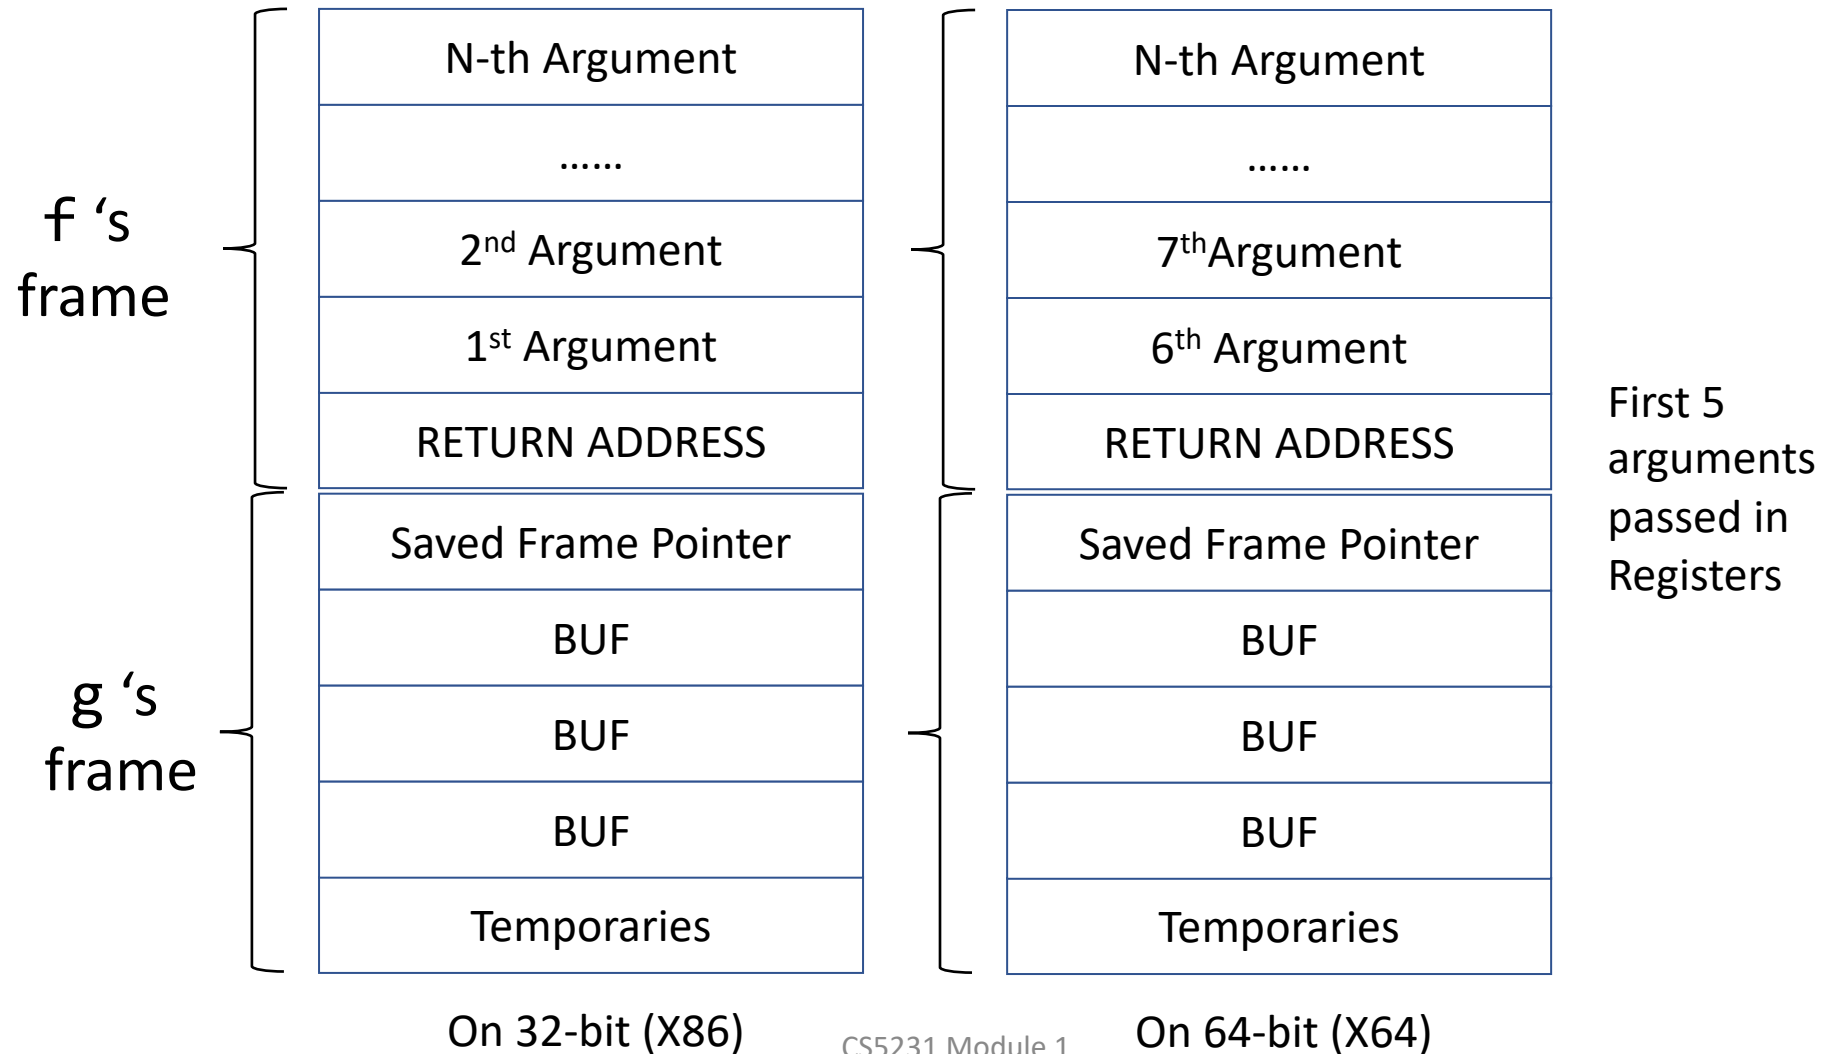
Spatial Memory Errors: Format String Bugs

Variable Argument Functions (Example – Printf / Scanf)



On 32-bit (X86)

Architectural Differences (x86 vs x64)



Format String Vulnerabilities

```
#include <stdio.h>

int main()
{
    srand(time(NULL));

    char localStr[100];
    int magicNumber = rand() % 100;
    int userCode = 0xBBBBBBBB;

    printf("Username? ");
    fgets(localStr, sizeof(localStr), stdin);

    printf("Hello ");
    printf(localStr);
    printf("What is the access code? ");

    scanf("%d", &userCode);
    if (userCode == magicNumber)
        printf("You win!\n");
    ...}
```

Format String Vulnerabilities

Format String

Argument 1 Argument 2

```
printf("Hello %s, you are %i years old", myName, myAge);
```

Malicious Format String

Stack.....

```
printf("AAAA %08x %08x %08x");
```

Format String Vulnerabilities

```
#include <stdio.h>
```

```
int main()
{
    srand(time(NULL));

    char localStr[100];
    int magicNumber = rand() % 100;
    int userCode = 0BBBBBBB;

    printf("Username? ");
    fgets(localStr, sizeof(localStr), stdin);
```


```
    printf("Hello ");
    printf(localStr);
```

AAAA %08x %08x %08x %08x %08x %08x %08x %08x %08x %08x %08x %08x %08x %08x %08x %08x %08x

```
    scanf("%d", &userCode);
    if (userCode == magicNumber)
        printf("You win!\n");
```

```
    ... }
```

The “main” stack frame



Format String Specifiers

Format String: Example “%s%d”

Format Specifiers

%d - print as number

%p - print as pointer

%c - print as character

%s - read from the address provided and print bytes until the NULL byte is reached

%n - write number of bytes already printed in the address provided

<n>\$ - accesses the nth positional argument with respect to printf (ex: %5\$p)

Temporal Memory Errors: Use-after-free & Double Free

Lifetime & Scope of Variables

- **Scope:**
 - Region of code where a variable can be accessed
 - E.g. Global, Function-local, Heap (dynamic)
 - **Lifetime:**
 - Portion of program execution during which storage is guaranteed
 - E.g. Auto vs. static
- Are Programming Language Abstractions
- Not instruction set / hardware abstractions

```
1. int z=0;
2. int g(int x, int y) {
3.     char* buf;
4.     buf = malloc (50);
5.     scanf("%s", buf);
6.     free (buf);
7. ...
8. }
```

Variable	Scope	Lifetime
z	Global, Line 1-8	Throughout the program
x	Local, Line 3-7	Execution of g
y	Local, Line 3-7	Execution of g
buf	Local, Line 3-7	Execution of g
"%s"	Constant Literal, Line 5	Execution of Line 5 (undefined?)
*buf	Heap, Line 4-7	Line 5-6

Question

```
int foo() {  
    int *p;  
    {  
        int x = 5;  
        p = &x;  
    }  
    return *p;  
}
```

What will this program return?

Answer:

Undefined behavior as per C11 standard

Why?


- The lifetime of x is within the inner {}
- The compiler can choose to remove the storage for “x”
- The pointer p is in scope at the last line
- The pointed-to object, however, is accessed out of scope!

Temporal Memory Errors

Example: Use-after-free

Temporal Mem Error: When program accesses mem. beyond its valid lifetime!

```
1 class Div: Element;
2 class Body: Element;
3 class Document {
4     Element* child;
5 };
6
7 // (a) memory allocations
8 Document *doc = new Document();
9 Body *body = new Body();
10 Div *div = new Div();
11
12 // (b) using memory: propagating pointers
13 doc->child = body;
14 body->child = div;
15
16 // (c) memory free: doc->child is now dangled
17 delete body;
18
```



Summary & Key Takeaways

- Memory Errors / Vulnerabilities
 - Spatial & Temporal
- Worst case: Can give attackers capability to read / write any value anywhere in memory
- Hardware does not give memory safety

